

Linux Foundation

CKA Exam

Certified Kubernetes Administrator (CKA) Program Exam

Product Questions: 122

Version: 4.0

Question: 1

CORRECT TEXT

Create a namespace called 'development' and a pod with image nginx called nginx on this namespace.

Answer:

```
kubectl create namespace development  
kubectl run nginx --image=nginx --restart=Never -n development
```

Question: 2

Create a nginx pod with label env=test in engineering namespace

See the solution below.

A. `kubectl run nginx --image=nginx --restart=Never --labels=env=test --namespace=engineering --dry-run -o yaml > nginx-pod.yaml`

`kubectl run nginx --image=nginx --restart=Never --labels=env=test --namespace=engineering --dry-run -o yaml | kubectl create -n engineering -f -`

YAML File:

```
apiVersion: v1  
kind: Pod  
metadata:  
name: nginx  
namespace: engineering  
labels:  
env: test  
spec:  
containers:  
- name: nginx  
image: nginx  
imagePullPolicy: IfNotPresent  
restartPolicy: Never
```

`kubectl create -f nginx-pod.yaml`

B. `kubectl run nginx --image=nginx --restart=Never --labels=env=test --namespace=engineering --dry-`

```
run -o yaml > nginx-pod.yaml
kubectl run nginx --image=nginx --restart=Never --labels=env=test --namespace=engineering --dry-run -o yaml | kubectl create -n engineering -f -
YAML File:
```

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
  namespace: engineering
  labels:
    env: test
spec:
  containers:
  - name: nginx
    image: nginx
    imagePullPolicy: IfNotPresent
    restartPolicy: Never
```

kubectl create -f nginx-pod.yaml

Answer: A

Question: 3

CORRECT TEXT

Get list of all pods in all namespaces and write it to file “/opt/pods-list.yaml”

Answer: .

kubectl get po –all-namespaces > /opt/pods-list.yaml

Question: 4

CORRECT TEXT

Create a pod with image nginx called nginx and allow traffic on port 80

Answer:

kubectl run nginx --image=nginx --restart=Never --port=80

Question: 5

CORRECT TEXT

Create a busybox pod that runs the command “env” and save the output to “envpod” file

Answer:

kubectl run busybox --image=busybox --restart=Never --rm -it -- env > envpod.yaml

Question: 6

CORRECT TEXT

List pod logs named “frontend” and search for the pattern “started” and write it to a file “/opt/error-logs”

Answer:

Kubectl logs frontend | grep -i “started” > /opt/error-logs

Question: 7

CORRECT TEXT

Create a pod that echo “hello world” and then exists. Have the pod deleted automatically when it’s completed

Answer:

```
kubectl run busybox --image=busybox -it --rm --restart=Never --  
/bin/sh -c 'echo hello world'  
kubectl get po # You shouldn't see pod with the name "busybox"
```

Question: 8

Create a pod with environment variables as var1=value1.Check the environment variable in pod

```
A. kubectl run nginx --image=nginx --restart=Never --env=var1=value1  
# then  
kubectl exec -it nginx -- env  
# or  
kubectl describe po nginx | grep value1  
B. kubectl run nginx --image=nginx --restart=Never --env=var1=value1  
# then  
kubectl exec -it nginx -- env  
# or  
kubectl exec -it nginx -- sh -c 'echo $var1'  
# or  
kubectl describe po nginx | grep value1
```

Answer: B

Question: 9

CORRECT TEXT

Get list of all the pods showing name and namespace with a jsonpath expression.

Answer:

```
kubectl get pods -o=jsonpath=".items[*]['metadata.name' , 'metadata.namespace']"
```

Question: 10

CORRECT TEXT

Check the image version in pod without the describe command

Answer:

```
kubectl get po nginx -o jsonpath='{.spec.containers[].image}{"\n"}
```

Question: 11

CORRECT TEXT

List the nginx pod with custom columns POD_NAME and POD_STATUS

Answer:

```
kubectl get po -o=custom-columns="POD_NAME:.metadata.name, POD_STATUS:.status.containerStatuses[].state"
```

Question: 12

CORRECT TEXT

List all the pods sorted by name

Answer:

```
kubect1 get pods --sort-by=.metadata.name
```

Question: 13

Create a pod that having 3 containers in it? (Multi-Container)

- A. image=nginx, image=redis, image=consul
- Name nginx container as “nginx-container”

Name redis container as “redis-container”
Name consul container as “consul-container”
Create a pod manifest file for a container and append container section for rest of the images
kubectl run multi-container --generator=run-pod/v1 --image=nginx --dry-run -o yaml > multi-container.yaml
then
vim multi-container.yaml
apiVersion: v1
kind: Pod
metadata:
labels:
run: multi-container
name: multi-container
spec:
containers:
- image: nginx
name: nginx-container
- image: redis
name: redis-container
- image: consul
name: consul-container
restartPolicy: Always

B. image=nginx, image=redis, image=consul
Name nginx container as “nginx-container”
Name redis container as “redis-container”
Name consul container as “consul-container”
Create a pod manifest file for a container and append container section for rest of the images
kubectl run multi-container --generator=run-pod/v1 --image=nginx --dry-run -o yaml > multi-container.yaml
then
vim multi-container.yaml
labels:
run: multi-container
name: multi-container
spec:
containers:
- image: nginx
name: nginx-container
- image: redis
name: consul-container
restartPolicy: Always

Answer: A

Question: 14

Create 2 nginx image pods in which one of them is labelled with env=prod and another one labelled with env=dev and verify the same.

A. `kubectl run --generator=run-pod/v1 --image=nginx -- labels=env=prod nginx-prod --dry-run -o yaml > nginx-prod-pod.yaml` Now, edit `nginx-prod-pod.yaml` file and remove entries like “`creationTimestamp: null`” “`dnsPolicy: ClusterFirst`”

`vim nginx-prod-pod.yaml`

`apiVersion: v1`

`kind: Pod`

`metadata:`

`labels:`

`env: prod`

`name: nginx-prod`

`spec:`

`containers:`

`- image: nginx`

`name: nginx-prod`

`restartPolicy: Always`

`# kubectl create -f nginx-prod-pod.yaml`

`kubectl run --generator=run-pod/v1 --image=nginx --`

`labels=env=dev nginx-dev --dry-run -o yaml > nginx-dev-pod.yaml`

`apiVersion: v1`

`kind: Pod`

`metadata:`

`labels:`

`env: dev`

`name: nginx-dev`

`spec:`

`containers:`

`- image: nginx`

`name: nginx-dev`

`restartPolicy: Always`

`# kubectl create -f nginx-dev-pod.yaml`

`Verify :`

`kubectl get po --show-labels`

`kubectl get po -l env=prod`

`kubectl get po -l env=dev`

B. `kubectl run --generator=run-pod/v1 --image=nginx -- labels=env=prod nginx-prod --dry-run -o yaml > nginx-prod-pod.yaml` Now, edit `nginx-prod-pod.yaml` file and remove entries like “`creationTimestamp: null`” “`dnsPolicy: ClusterFirst`”

`vim nginx-prod-pod.yaml`

`apiVersion: v1`

`kind: Pod`

```

metadata:
labels:
env: prod
name: nginx-prod
spec:
containers:
- image: nginx
name: nginx-prod
restartPolicy: Always
# kubectl create -f nginx-prod-pod.yaml
kubectl run --generator=run-pod/v1 --image=nginx --
labels=env=dev nginx-dev --dry-run -o yaml > nginx-dev-pod.yaml
apiVersion: v1
kind: Pod
metadata:
- image: nginx
name: nginx-dev
restartPolicy: Always
# kubectl create -f nginx-prod-dev.yaml
Verify :
kubectl get po --show-labels
kubectl get po -l env=dev

```

Answer: A

Question: 15

Get IP address of the pod – “nginx-dev”

- A. Kubectl get po -o wide
Using JsonPath
kubectl get pods
.items[*]{.metadata.name}{"\t"}{.status.podIP}{"\n"}{end}'
B. Kubectl get po -o wide
Using JsonPath
kubectl get pods -o=jsonpath='{range
.items[*]}{.metadata.name}{"\t"}{.status.podIP}{"\n"}{end}'

Answer: B

Question: 16

CORRECT TEXT

Print pod name and start time to “/opt/pod-status” file

Answer:

```

kubectl get pods -o=jsonpath='{range
.items[*]}{.metadata.name}{"\t"}{.status.podIP}{"\n"}{end}'

```

Question: 17

CORRECT TEXT

Check the Image version of nginx-dev pod using jsonpath

Answer:

```
kubectl get po nginx-dev -o  
jsonpath='{.spec.containers[].image}{"\n"}'
```

Question: 18

CORRECT TEXT

Create a busybox pod and add “sleep 3600” command

Answer:

```
kubectl run busybox --image=busybox --restart=Never -- /bin/sh -c  
"sleep 3600"
```

Question: 19

Create an nginx pod and list the pod with different levels of verbosity

A. // create a pod

```
kubectl run nginx --image=nginx --restart=Never --port=80
```

// List the pod with different verbosity

```
kubectl get po nginx --v=7
```

```
kubectl get po nginx --v=8
```

```
kubectl get po nginx --v=9
```

B. // create a pod

```
kubectl run nginx --image=nginx --restart=Never --port=80
```

// List the pod with different verbosity

```
kubectl get po nginx --v=7
```

```
kubectl get po nginx --v=6
```

```
kubectl get po nginx --v=9
```

Answer: A

Question: 20

CORRECT TEXT

List the nginx pod with custom columns POD_NAME and POD_STATUS

Answer:

```
kubectl get po -o=custom-columns="POD_NAME:.metadata.name,  
POD_STATUS:.status.containerStatuses[].state"
```

Question: 21

CORRECT TEXT

List all the pods sorted by name

```
kubectl get pods --sort-by=.metadata.name
```

Answer:

Question: 22

CORRECT TEXT

List all the pods sorted by created timestamp

```
kubect1 get pods--sort-by=.metadata.creationTimestamp
```

Answer:

Question: 23

CORRECT TEXT

List all the pods showing name and namespace with a json path expression

Answer:

```
kubectl get pods -o=jsonpath=".items[*]['metadata.name',  
'metadata.namespace']"
```

Question: 24

List “nginx-dev” and “nginx-prod” pod and delete those pods

A. kubect1 get pods -o wide
kubectl delete po
kubectl delete po “nginx-prod”

B. kubect1 get pods -o wide
kubectl delete po “nginx-dev”
kubectl delete po “nginx-prod”

Answer: B

Question: 25

CORRECT TEXT

Delete the pod without any delay (force delete)

Answer:

Kubectl delete po "POD-NAME" --grace-period=0 -force

Question: 26

Create a redis pod and expose it on port 6379

A. kubectl run redis --image=redis --restart=Never --port=6379

YAML File :

```
apiVersion: v1
kind: Pod
metadata:
labels:
run: redis
name: redis
spec:
containers:
- image: redis
name: redis
ports:
- containerPort: 6379
Rt restartPolicy: Always
```

B. kubectl run redis --image=redis --restart=Never --port=6379

YAML File :

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
labels:
```

```
run: redis
```

```
name: redis
```

```
spec:
```

```
containers:
```

```
ports:
```

```
- containerPort: 6679
```

```
Rt restartPolicy: Alwaysf
```

Answer: A

Question: 27

CORRECT TEXT

Create the nginx pod with version 1.17.4 and expose it on port 80

Answer:

```
kubectl run nginx --image=nginx:1.17.4 --restart=Never --  
port=80
```

Question: 28

Change the Image version to 1.15-alpine for the pod you just created and verify the image version is updated.

A. Kubectl set image pod/nginx nginx=nginx:1.15-alpine
kubectl describe po nginx
// another way it will open vi editor and change the version
kubectl edit po nginx
kubectl describe po nginx

B. Kubectl set image pod/nginx nginx=nginx:1.15-alpine
kubectl describe po nginx
// another way it will open vi editor and change the version
kubectl describe po nginx

Answer: A

Question: 29

CORRECT TEXT

Change the Image version back to 1.17.1 for the pod you just updated and observe the changes

Answer:

```
kubectl set image pod/nginx nginx=nginx:1.17.1  
kubectl describe po nginx  
kubectl get po nginx -w # watch it
```

Question: 30

Create a redis pod, and have it use a non-persistent storage

Note: In exam, you will have access to kubernetes.io site,

Refer : <https://kubernetes.io/docs/tasks/configure-pod-container/configurevolume-storage/>

A. apiVersion: v1
kind: Pod
metadata:
name: redis
spec:
containers:
- name: redis
image: redis
volumeMounts:
- name: redis-storage
mountPath: /data/redis
ports:
- containerPort: 6379

```
volumes:  
- name: redis-storage  
emptyDir: {}  
B. apiVersion: v1  
kind: Pod  
metadata:  
name: redis  
spec:  
containers:  
- name: redis  
image: redis  
volumeMounts:  
- containerPort: 6379  
volumes:  
- name: redis-storage  
emptyDir: {}
```

Answer: A

Question: 31

Create a Pod with three busy box containers with commands “ls; sleep 3600;”, “echo Hello World; sleep 3600;” and “echo this is the third container; sleep 3600” respectively and check the status

A. // first create single container pod with dry run flag
kubectl run busybox --image=busybox --restart=Always --dry-run
-o yaml -- bin/sh -c "sleep 3600; ls" > multi-container.yaml
// edit the pod to following yaml and create it

```
apiVersion: v1  
kind: Pod  
metadata:  
labels:  
run: busybox  
name: busybox  
spec:  
containers:  
- args:  
- bin/sh  
- -C  
- ls; sleep 3600  
- echo Hello world; sleep 3600  
image: busybox  
name: busybox-container-2  
- args:  
- bin/sh  
- -C  
- echo this is third container; sleep 3600  
image: busybox
```

```
name: busybox-container-3
restartPolicy: Always
// Verify
Kubectl get pods
B. // first create single container pod with dry run flag
kubectl run busybox --image=busybox --restart=Always --dry-run
-o yaml -- bin/sh -c "sleep 3600; ls" > multi-container.yaml
// edit the pod to following yaml and create it
apiVersion: v1
kind: Pod
metadata:
labels:
run: busybox
name: busybox
spec:
containers:
- args:
- bin/sh
- -c
- ls; sleep 3600
image: busybox
name: busybox-container-1
- args:
- bin/sh
- -c
- echo Hello world; sleep 3600
image: busybox
name: busybox-container-2
- args:
- bin/sh
- -c
- echo this is third container; sleep 3600
image: busybox
name: busybox-container-3
restartPolicy: Always
// Verify
Kubectl get pods
```

Answer: B

Question: 32

Check logs of each container that “busyboxpod-{1,2,3}”

- A. kubectl logs busybox -c busybox-container-1
- kubectl logs busybox -c busybox-container-2
- kubectl logs busybox -c busybox-container-3
- B. kubectl logs busybox -c busybox-container-1

```
kubectl logs busybox -c busybox-container-3  
kubectl logs busybox -c busybox-container-3
```

Answer: A

Question: 33

Create a Pod with main container busybox and which executes this “while true; do echo ‘Hi I am from Main container’ >> /var/log/index.html; sleep 5; done” and with sidcar container with nginx image which exposes on port 80. Use emptyDir Volume and mount this volume on path /var/log for busybox and on path /usr/share/nginx/html for nginx container. Verify both containers are running.

A. // create an initial yaml file with this

```
kubectl run multi-cont-pod --image=busbox --restart=Never --  
dry-run -o yaml > multi-container.yaml  
// edit the yml as below and create it  
kubectl create -f multi-container.yaml  
vim multi-container.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
labels:  
run: multi-cont-pod  
name: multi-cont-pod  
spec:  
volumes:  
- name: var-logs  
emptyDir: {}  
containers:  
- image: busybox  
command: ["/bin/sh"]  
args: ["-c", "while true; do echo 'Hi I am from Main  
container' >> /var/log/index.html; sleep 5;done"]  
name: main-container  
volumeMounts:  
- name: var-logs  
mountPath: /var/log  
- image: nginx  
name: sidcar-container  
ports:  
- containerPort: 80  
volumeMounts:  
- name: var-logs  
mountPath: /usr/share/nginx/html  
restartPolicy: Never
```

```
// Create Pod
kubectl apply -f multi-container.yaml
//Verify
kubectl get pods
B. // create an initial yaml file with this
kubectl run multi-cont-pod --image=busybox --restart=Never --
dry-run -o yaml > multi-container.yaml
// edit the yml as below and create it
kubectl create -f multi-container.yaml
vim multi-container.yaml
apiVersion: v1
kind: Pod
metadata:
labels:
run: multi-cont-pod
name: multi-cont-pod
spec:
volumes:
- image: busybox
command: ["/bin/sh"]
args: ["-c", "while true; do echo 'Hi I am from Main
container' >> /var/log/index.html; sleep 5;done"]
name: main-container
volumeMounts:
- name: var-logs
mountPath: /var/log
- image: nginx
name: sidecar-container
ports:
mountPath: /usr/share/nginx/html
restartPolicy: Never
// Create Pod
kubectl apply -f multi-container.yaml
//Verify
kubectl get pods
```

Answer: A

Question: 34

CORRECT TEXT

Exec into both containers and verify that main.txt exist and

NEED TO WRITE ANSWER FOR THIS

Answer:

Question: 35

Create an nginx pod and set an env value as 'var1=val1'. Check the env value existence within the pod

A. `kubectl run nginx --image=nginx --restart=Never --env=var1=val1`
then
`kubectl exec -it nginx -- env`
or
`kubectl exec -it nginx -- sh -c 'echo $var1'`
or
`kubectl describe po nginx | grep val1`
or
`kubectl run nginx --restart=Never --image=nginx --env=var1=val1`
-it --rm -- env
B. `kubectl run nginx --image=nginx --restart=Never --env=var1=val1`
then
`kubectl exec -it nginx -- env`
or
`kubectl run nginx --restart=Never --image=nginx --env=var1=val1`
-it --rm -- env

Answer: A

Question: 36

Create a pod with init container which create a file "test.txt" in "workdir" directory. Main container should check a file "test.txt" exists and execute sleep 9999 if the file exists.

A. // create an initial yaml file with this
`kubectl run init-cont-pod --image=alpine --restart=Never --dry-run -o yaml > init-cont-pod.yaml`
// edit the yml as below and create it
`vim init-cont-pod.yaml`
`apiVersion: v1`
`kind: Pod`
`metadata:`
`name: init-cont-pod`
`labels:`
`app: myapp`
`spec:`
`volumes:`
`- name: test-volume`
`emptyDir: {}`
`containers:`

```
- name: main-container
image: busybox:1.28
command: ['sh', '-c', 'if [ -f /workdir/test.txt ]; then sleep
9999; fi']
volumeMounts:
- name: test-volume
mountPath: /workdir
initContainers:
- name: init-myservice
image: busybox:1.28
command: ['sh', '-c', "mkdir /workdir; echo >
/workdir/test.txt"]
volumeMounts:
- name: test-volume
mountPath: /workdir
// Create the pod
kubectl apply -f init-cont-pod.yaml
kubectl get pods
// Check Events by doing
kubectl describe po init-cont-pod
Init Containers:
init-myservice:
Container ID:
docker://ebdbf5fad1c95111d9b0e0e2e743c2e347c81b8d4eb5abcccdfe1dd74524
0d4f
Image: busybox:1.28
Image ID: dockerpullable://busybox@sha256:141c253bc4c3fd0a201d32dc1f493bcf3fff003b6df
416dea4f41046e0f37d47
Port: <none>
Host Port: <none>
Command:
sh
-c
mkdir /workdir; echo > /workdir/test.txt
State: Terminated
Reason: Completed
B. // create an initial yaml file with this
kubectl run init-cont-pod --image=alpine --restart=Never --dry-run -o
yaml > init-cont-pod.yaml
// edit the yaml as below and create it
vim init-cont-pod.yaml
apiVersion: v1
kind: Pod
metadata:
name: init-cont-pod
labels:
app: myapp
spec:
```

```

volumes:
- name: test-volume
emptyDir: {}
containers:
- name: main-container
image: busybox:1.28
command: ['sh', '-c', 'if [ -f /workdir/test.txt ]; then sleep
9999; fi']
volumeMounts:
image: busybox:1.28
command: ['sh', '-c', "mkdir /workdir; echo >
/workdir/test.txt"]
volumeMounts:
- name: test-volume
mountPath: /workdir
// Create the pod
kubectl apply -f init-cont-pod.yaml
kubectl get pods
// Check Events by doing
kubectl describe po init-cont-pod
Init Containers:
init-myservice:
Container ID:
docker://ebdbf5fad1c95111d9b0e0e2e743c2e347c81b8d4eb5abcccdfe1dd74524
0d4f
Image: busybox:1.28
Image ID: dockerpullable://busybox@sha256:141c253bc4c3fd0a201d32dc1f493bcf3fff003b6df
416dea4f41046e0f37d47
Port: <none>
Host Port: <none>
Command:
sh
-c
mkdir /workdir; echo > /workdir/test.txt
State: Terminated
Reason: Completed

```

Answer: A

Question: 37

Create a pod with init container which waits for a service called “myservice” to be created. Once init container completes, the myapp-container should start and print a message “The app is running” and sleep for 3600 seconds.

A. vim multi-container-pod.yaml

apiVersion: v1

kind: Pod

```
metadata:  
name: myapp-pod  
labels:  
app: myapp  
spec:  
containers:  
- name: myapp-container  
image: busybox:1.28  
command: ['sh', '-c', 'echo The app is running! && sleep  
3600']  
initContainers:  
- name: init-myservice  
done"]  
// Check whether service called "myservice" exists  
kubectl get svc  
Note: Pod will not start if service called "myservice" doesn't  
exist.  
// Now, Create the pod  
kubectl apply -f multi-container-pod.yaml  
B. vim multi-container-pod.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
name: myapp-pod  
labels:  
app: myapp  
spec:  
containers:  
- name: myapp-container  
image: busybox:1.28  
command: ['sh', '-c', 'echo The app is running! && sleep  
3600']  
initContainers:  
- name: init-myservice  
image: busybox:1.28  
command: ['sh', '-c', "until nslookup myservice.$(cat  
/var/run/secrets/kubernetes.io/serviceaccount/namespace).s  
vc.cluster.local; do echo waiting for myservice; sleep 2;  
done"]  
// Check whether service called "myservice" exists  
kubectl get svc  
Note: Pod will not start if service called "myservice" doesn't  
exist.  
// Now, Create the pod  
kubectl apply -f multi-container-pod.yaml
```

Answer: B

Question: 38

Create 5 nginx pods in which two of them is labeled env=prod and three of them is labeled env=dev

A. kubectl run nginx-dev1 --image=nginx --restart=Never --
labels=env=dev
kubectl run nginx-dev2 --image=nginx --restart=Never --
labels=env=dev
kubectl run nginx-dev3 --image=nginx --restart=Never --
labels=env=dev
kubectl run nginx-prod1 --image=nginx --restart=Never --
labels=env=prod
kubectl run nginx-prod2 --image=nginx --restart=Never --
labels=env=prod
B. kubectl run nginx-dev1 --image=nginx --restart=Never --
labels=env=dev
kubectl run nginx-dev2 --image=nginx --restart=Never --
labels=env=dev
kubectl run nginx-prod1 --image=nginx --restart=Never --
labels=env=prod
kubectl run nginx-prod2 --image=nginx --restart=Never --
labels=env=prod

Answer: A

Question: 39

CORRECT TEXT

Get the pods with label env=dev and output the labels

kubectl get pods -l env=dev --show-labels

Answer:

Question: 40

CORRECT TEXT

Get the pods with labels env=dev and env=prod and output the labels as well

kubectl get pods -l 'env in (dev,prod)' --show-labels

Answer:

Question: 41

CORRECT TEXT

Get all the pods with label "env"

kubectl get pods -L env

Answer:

Question: 42

CORRECT TEXT

Change the label for one of the pod to env=uat and list all the pods to verify

Answer:

```
kubectl label pod/nginx-dev3 env=uat --overwrite  
kubectl get pods --show-labels
```

Question: 43

CORRECT TEXT

Get list of all the nodes with labels

Answer:

```
kubectl get nodes --show-labels
```

Question: 44

Create a nginx pod that will be deployed to node with the label "gpu=true"

```
A. kubectl run nginx --image=nginx --restart=Always --dry-run -o  
yaml > nodeselector-pod.yaml  
// add the nodeSelector like below and create the pod  
kubectl apply -f nodeselector-pod.yaml  
vim nodeselector-pod.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
name: nginx  
spec:  
nodeSelector:  
gpu: true  
containers:  
- image: nginx  
name: nginx  
restartPolicy: Always  
kubectl apply -f nodeselector-pod.yaml  
//Verify  
kubectl get no --show-labels  
kubectl get po  
kubectl describe po nginx | grep Node-Selectors  
B. kubectl run nginx --image=nginx --restart=Always --dry-run -o
```

```

yaml > nodeselector-pod.yaml
// add the nodeSelector like below and create the pod
kubectl apply -f nodeselector-pod.yaml
vim nodeselector-pod.yaml
apiVersion: v1
kind: Pod
metadata:
name: nginx
spec:
nodeSelector:
gpu: true
yaml
//Verify
kubectl get no --show-labels
kubectl get po
kubectl describe po nginx | grep Node-Selectors

```

Answer: A

Question: 45

Annotate the pod with name=webapp

A. kubectl annotate pod nginx-dev-pod name=webapp
 kubectl annotate pod nginx-prod-pod name=webapp
 // Verify
 kubectl describe po nginx-dev-pod | grep -i annotations
 B. kubectl annotate pod nginx-dev-pod name=webapp
 kubectl annotate pod nginx-prod-pod name=webapp
 // Verify
 kubectl describe po nginx-dev-pod | grep -i annotations
 kubectl describe po nginx-prod-pod | grep -i annotations

Answer: B

Question: 46

Create a deployment called webapp with image nginx having 5 replicas in it, put the file in /tmp directory with named webapp.yaml

A. //Create a file using dry run command
 kubectl create deploy --image=nginx --dry-run -o yaml >
 /tmp/webapp.yaml
 // Now, edit file webapp.yaml and update replicas=5
 apiVersion: apps/v1
 kind: Deployment
 metadata:

```
labels:  
app: webapp  
name: webapp  
spec:  
replicas: 5  
selector:  
matchLabels:  
app: webapp  
template:  
metadata:  
labels:  
app: webapp  
spec:  
containers:  
- image: nginx  
name: nginx  
Note: Search "deployment" in kubernetes.io site , you will get  
the page  
https://kubernetes.io/docs/concepts/workloads/controllers/deployment/  
// Verify the Deployment  
kubectl get deploy webapp --show-labels  
// Output the YAML file of the deployment webapp  
kubectl get deploy webapp -o yaml  
B. //Create a file using dry run command  
kubectl create deploy --image=nginx --dry-run -o yaml >  
/tmp/webapp.yaml  
// Now, edit file webapp.yaml and update replicas=5  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
labels:  
app: webapp  
name: webapp  
spec:  
replicas: 5  
selector:  
matchLabels:  
app: webapp  
template:  
metadata:  
labels:  
Note: Search "deployment" in kubernetes.io site , you will get  
the page  
https://kubernetes.io/docs/concepts/workloads/controllers/deployment/  
// Verify the Deployment  
kubectl get deploy webapp --show-labels
```

```
// Output the YAML file of the deployment webapp  
kubectl get deploy webapp -o yaml
```

Answer: A

Question: 47

Get the list of pods of webapp deployment

- A. // Get the label of the deployment
kubectl get deploy --show-labels
kubectl get pods -l app=webapp
- B. // Get the label of the deployment
kubectl get deploy --show-labels
// Get the pods with that label
kubectl get pods -l app=webapp

Answer: B

Question: 48

CORRECT TEXT

Scale the deployment from 5 replicas to 20 replicas and verify

```
kubectl scale deploy webapp --replicas=20  
kubectl get deploy webapp  
kubectl get po -l app=webapp
```

Answer:

Question: 49

CORRECT TEXT

Get the deployment rollout status

```
kubectl rollout status deploy webapp
```

Answer:

Question: 50

Deployment

- a. Create a deployment of webapp with image nginx:1.17.1 with container port 80 and verify the image version

A. // Create initial YAML file with --dry-run option

```
kubectl create deploy webapp --image=nginx:1.17.1 --dryrun=client -o yaml > webapp.yaml  
vim webapp.yaml  
apiVersion: apps/v1  
kind: Deployment
```

```
metadata:  
labels:  
app: webapp  
name: webapp  
spec:  
replicas: 1  
selector:  
matchLabels:  
app: webapp  
template:  
metadata:  
labels:  
app: webapp  
spec:  
containers:  
- image: nginx:1.17.1  
name: nginx  
kubectl create -f webapp.yaml --record=true  
//Verify Image Version  
kubectl describe deploy webapp | grep -i "Image"  
Using JsonPath  
kubectl get deploy -o=jsonpath='{range.items [*]}{[*]}  
.metadata.name{"\t"}{.spec.template.spec.containers[*].i  
mage}"\n}'  
B. // Create initial YAML file with --dry-run option  
kubectl create deploy webapp --image=nginx:1.17.1 --dryrun=client -o yaml > webapp.yaml  
vim webapp.yaml  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
labels:  
app: webapp  
name: webapp  
spec:  
replicas: 1  
containers:  
- image: nginx:1.17.1  
name: nginx  
kubectl create -f webapp.yaml --record=true  
//Verify Image Version  
kubectl describe deploy webapp | grep -i "Image"  
Using JsonPath  
kubectl get deploy -o=jsonpath='{range.items [*]}{[*]}  
.metadata.name{"\t"}{.spec.template.spec.containers[*].i  
mage}"\n}'
```

Answer: A

Question: 51

Update the deployment with the image version 1.17.4 and verify

A. kubectl set image deploy/webapp nginx=nginx:1.17.4
//Verify
kubectl describe deploy webapp | grep Image
kubectl get deploy -o=jsonpath='{range.items [*]}{[*]}\n{.metadata.name}{"\t"}{.spec.template.spec.containers[*].image}{"\n"}'
B. kubectl set image deploy/webapp nginx=nginx:1.17.4
//Verify
kubectl describe deploy webapp | grep Image
kubectl get deploy -
{.metadata.name}{"\t"}{.spec.template.spec.containers[*].image}{"\n"}

Answer: A

Question: 52

Check the rollout history and make sure everything is ok after the update

A. kubectl rollout history deploy webapp
kubectl get deploy webapp --show-labels
kubectl get rs -l app=webapp
kubectl get po -l app=webapp
B. kubectl rollout history deploy webapp
kubectl get deploy webapp --show-labels
kubectl get rs -
kubectl get po -l app=webapp

Answer: A

Question: 53

CORRECT TEXT

Undo the deployment to the previous version 1.17.1 and verify Image has the previous version

Answer:

kubectl rollout undo deploy webapp
kubectl describe deploy webapp | grep Image

Question: 54

CORRECT TEXT

Update the deployment with the image version 1.16.1 and verify the image and check the rollout history

Answer:

```
kubectl set image deploy/webapp nginx=nginx:1.16.1  
kubectl describe deploy webapp | grep Image  
kubectl rollout history deploy webapp
```

Question: 55

CORRECT TEXT

Check the history of deployment

Answer:

```
kubectl rollout history deployment webapp
```

Question: 56

CORRECT TEXT

Undo the deployment with the previous version and verify everything is Ok

Answer:

```
kubectl rollout undo deploy webapp  
kubectl rollout status deploy webapp  
kubectl get pods
```

Question: 57

Undo/Rollback deployment to specific revision “1”

A. // Check Deployment History
kubectl rollout history deployment webapp
//Rollback to particular revision
kubectl rollout undo deployment webapp --to-revision=1

B. // Check Deployment History
kubectl rollout history deployment webapp
kubectl rollout undo deployment webapp --to-revision=1

Answer: A

Question: 58

CORRECT TEXT

Check the history of the specific revision of that deployment

Answer:

```
kubectl rollout history deploy webapp --revision=3
```

Question: 59

CORRECT TEXT

Pause the rollout of the deployment

kubectl rollout pause deploy webapp

Answer:

Question: 60

CORRECT TEXT

Resume the rollout of the deployment

kubectl rollout resume deploy webapp

Answer:

Question: 61

CORRECT TEXT

Scale the deployment to 5 replicas

kubectl scale deployment webapp --replicas=5
//Verify
kubectl get deploy
kubectl get po,rs

Answer:

Question: 62

CORRECT TEXT

Scale down the deployment to 1 replica

kubectl scale deployment webapp --replicas=1
//Verify
kubectl get deploy
kubectl get po,rs

Answer:

Question: 63

CORRECT TEXT

Apply the autoscaling to this deployment with minimum 10 and maximum 20 replicas and target CPU of 85% and verify hpa is created and replicas are increased to 10 from 1

Answer:

```
kubectl autoscale deploy webapp --min=10 --max=20 --cpu  
percent=85  
kubectl get hpa
```

```
kubectl get pod -l app=webapp
```

Question: 64

CORRECT TEXT

Clean the cluster by deleting deployment and hpa you just created

```
kubectl delete deploy webapp  
kubectl delete hpa webapp
```

Answer:

Question: 65

Create a deployment named “myapp” that having 2 replicas with nginx image and expose deployment as service named “myservice”

A. // Create a YAML Template

```
kubectl create deploy myapp --image=nginx --dry-run -o yaml >  
myapp.yaml  
//Update replicas=2 in myapp.yaml file  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
labels:  
app: myapp  
name: myapp  
spec:  
replicas: 2  
selector:  
matchLabels:  
app: myapp  
template:  
metadata:  
labels:  
app: myapp  
spec:  
containers:  
- image: nginx  
name: nginx  
// Create deployment  
kubectl create -f myapp.yaml  
// Creating YAML template for service
```

```
kubectl expose deployment myapp --type=ClusterIP --port=80 --  
target-port=80 --name=myservice --dry-run -o yaml >  
myservice.yaml  
YAML File:  
apiVersion: v1  
kind: Service  
metadata:  
labels:  
app: myapp  
name: myservice  
spec:  
ports:  
- port: 80  
protocol: TCP  
targetPort: 80  
selector:  
app: myapp  
type: ClusterIP  
kubectl get svc  
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S)  
AGE  
kubernetes ClusterIP 10.2.0.1 <none> 443/TCP  
158d  
myservice ClusterIP 10.2.96.175 <none> 80/TCP  
40s  
B. // Create a YAML Template  
kubectl create deploy myapp --image=nginx --dry-run -o yaml >  
myapp.yaml  
//Update replicas=2 in myapp.yaml file  
apiVersion: apps/v1  
kind: Deployment  
metadata:  
labels:  
app: myapp  
name: myapp  
spec:  
replicas: 2  
selector:  
matchLabels:  
app: myapp  
template:  
metadata:  
labels:  
app: myapp  
spec:  
containers:  
- image: nginx  
name: nginx
```

```
// Create deployment
kubectl create -f myapp.yaml
// Creating YAML template for service
kubectl expose deployment myapp --type=ClusterIP --port=60 --
target-port=60 --name=myservice --dry-run -o yaml >
myservice.yaml
YAML File:
apiVersion: v1
kind: Service
metadata:
labels:
app: myapp
name: myservice
spec:
ports:
- port: 60
protocol: TCP
targetPort: 80
selector:
app: myapp
type: ClusterIP
kubectl get svc
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S)
AGE
kubernetes ClusterIP 10.2.0.1 <none> 443/TCP
158d
myservice ClusterIP 10.2.96.175 <none> 80/TCP
40s
```

Answer: A

Question: 66

Create the deployment redis with image=redis and expose it with “NodePort” service redis-service

A. kubectl create deploy redis --image=redis --dry-run -o yaml >
redis-deploy.yaml
Edit redis-deploy.yaml file
vim redis-deploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
labels:
app: redis
name: redis
spec:
replicas: 1
selector:

```
matchLabels:  
app: redis  
template:  
metadata:  
labels:  
app: redis  
spec:  
containers:  
- image: redis  
name: redis  
//Creating Service  
kubectl expose deploy redis --type=NodePort --port=6379 --  
target-port=6379 --name redis-service  
// Verify  
kubectl get deploy,svc  
B. kubectl create deploy redis --image=redis --dry-run -o yaml >  
redis-deploy.yaml  
Edit redis-deploy.yaml file  
name: redis  
spec:  
replicas: 1  
selector:  
matchLabels:  
app: redis  
template:  
metadata:  
labels:  
app: redis  
spec:  
containers:  
- image: redis  
name: redis  
//Creating Service  
kubectl expose deploy redis --type=NodePort --port=6379 --  
target-port=6379 --name redis-service  
// Verify  
kubectl get deploy,svc
```

Answer: A

Question: 67

Get the DNS records for the service and pods for the deployment redis and the put the value in /tmp/dnsrecordpod and /tmp/dnsrecord-service

A. // Get Pod Ip
kubectl get po -o wide
// Get Service Name

```
kubectl get svc
// Create a temporary pod and execute nslookup command
Note: POD IP format should be a-b-c-d and not a.b.c.d
kubectl run busybox --image=busybox:1.28 --restart=Never -
-rm -it -- nslookup 192-168-0-69.default.pod >
/tmp/dnsrecord-pod
kubectl run busybox1 --image=busybox:1.26 --restart=Never
--rm -it -- nslookup redis-service > /tmp/dnsrecordservice
//Verify
cat /tmp/dnsrecord-pod
Server: 10.2.8.10
Address 1: 10.2.0.10 kube-dns.kube
system.svc.cluster.local
Name: 192-168-0-69.default.pod
Address 1: 192.168.0.69 192-168-0-69.redis
service.default.svc.cluster.local
cat /tmp/dnsrecord-pod
Server: 10.2.0.10
Address 1: 10.2.0.10 kube-dns.kube
system.svc.cluster.local
Name: 192-168-0-69.default.pod
Address 1: 192.168.0.69 192-168-0-69.redis
service.default.svc.cluster.local
B. // Get Pod Ip
kubectl get po -o wide
// Get Service Name
kubectl get svc
// Create a temporary pod and execute nslookup command
Note: POD IP format should be a-b-c-d and not a.b.c.d
kubectl run busybox --image=busybox:1.28 --restart=Never -
-rm -it -- nslookup 192-168-0-69.default.pod >
/tmp/dnsrecord-pod
kubectl run busybox1 --image=busybox:1.28 --restart=Never
--rm -it -- nslookup redis-service > /tmp/dnsrecordservice
//Verify
cat /tmp/dnsrecord-pod
Server: 10.2.0.10
Address 1: 10.2.0.10 kube-dns.kube
system.svc.cluster.local
Name: 192-168-0-69.default.pod
Address 1: 192.168.0.69 192-168-0-69.redis
service.default.svc.cluster.local
cat /tmp/dnsrecord-pod
Server: 10.2.0.10
Address 1: 10.2.0.10 kube-dns.kube
system.svc.cluster.local
Name: 192-168-0-69.default.pod
Address 1: 192.168.0.69 192-168-0-69.redis
```

service.default.svc.cluster.local

Answer: B

Question: 68

CORRECT TEXT

List all the pods that are serviced by the service “webservice” and copy the output in /opt/\$USER/webservice.targets

Note: You need to list the endpoints

Answer:

```
kubectl describe svc webservice | grep -i "Endpoints" >
/opt/$USER/webservice.targets
kubectl get endpoints webservice >
/opt/$USER/webservice.targets
```

Question: 69

Create a Job with an image node which prints node version and verifies there is a pod created for this job

A. kubectl create job nodeversion --image=node -- node -v
kubectl get job -w
kubectl get pod
YAML File:
apiVersion: batch/v1
kind: Job
metadata:
labels:
job-name: nodeversion
name: nodeversion
spec:
completions: 1
parallelism: 1
selector:
matchLabels:
job-name: nodeversion
template:
metadata:
labels:
job-name: nodeversion
spec:
containers:
- command:
- node
- -v
image: node

```
imagePullPolicy: Always
name: nodeversion
restartPolicy: Never
B. kubectl create job nodeversion --image=node -- node -v
kubectl get job -w
kubectl get pod
YAML File:
apiVersion: batch/v1
kind: Job
metadata:
labels:
job-name: nodeversion
name: nodeversion
spec:
completions: 1
parallelism: 1
labels:
job-name: nodeversion
spec:
containers:
- command:
- node
- -v
image: node
imagePullPolicy: Always
name: nodeversion
restartPolicy: Never
```

Answer: A

Question: 70

Create a job named “hello-job” with the image busybox which echos “Hello I’m running job”

```
A. kubectl create job hello-job --image=busybox --dry-run -o yaml
-- echo "Hello I'm running job" > hello-job.yaml
kubectl create -f hello-job.yaml
//Verify Job
kubectl get job
kubectl get po
kubectl logs hello-job-*
B. kubectl create job hello-job --image=busybox --dry-run -o yaml
-- echo "Hello I'm running job" > hello-job.yaml
kubectl create -f hello-job.yaml
//Verify Job
kubectl get po
kubectl logs hello-job-*
```

Answer: A

Question: 71

Modify “hello-job” and make it run 10 times one after one and 5 times parallelism: 5

A. kubectl create job hello-job --image=busybox --dry-run -o yaml
-- echo "Hello I am from job" > hello-job.yaml
// edit the yaml file to add completions: 16 and
kubectl create -f hello-job.yaml

YAML File:

```
apiVersion: batch/v1
```

```
kind: Job
```

```
metadata:
```

```
name: hello-job
```

```
spec:
```

```
completions: 16
```

```
parallelism: 5
```

```
template:
```

```
metadata:
```

```
spec:
```

```
containers:
```

```
- command:
```

```
- echo
```

```
- Hello I am from job
```

```
image: busybox
```

```
name: hello-job
```

```
restartPolicy: Never
```

B. kubectl create job hello-job --image=busybox --dry-run -o yaml

-- echo "Hello I am from job" > hello-job.yaml

// edit the yaml file to add completions: 10 and

kubectl create -f hello-job.yaml

YAML File:

```
apiVersion: batch/v1
```

```
kind: Job
```

```
metadata:
```

```
name: hello-job
```

```
spec:
```

```
completions: 10
```

```
parallelism: 5
```

```
template:
```

```
metadata:
```

```
spec:
```

```
containers:
```

```
- command:
```

```
- echo
```

```
- Hello I am from job
```

```
image: busybox
```

```
name: hello-job
restartPolicy: Never
```

Answer: B

Question: 72

CORRECT TEXT

Watch the job that runs 10 times one by one and verify 10 pods are created and delete those after it's completed

Answer:

```
kubectl get job -w
kubectl get po
kubectl delete job hello-job
```

Question: 73

Create a Cronjob with busybox image that prints date and hello from kubernetes cluster message for every minute

A. CronJob Syntax:

```
* --> Minute
* --> Hours
* --> Day of The Month
* --> Month
* --> Day of the Week
*/1 * * * * --> Execute a command every one minutes.
vim date-job.yaml
apiVersion: batch/v1beta1
kind: CronJob
metadata:
  name: date-job
spec:
  schedule: "*/1 * * * *"
  jobTemplate:
    spec:
      template:
        spec:
          containers:
            - name: hello
              image: busybox
              args:
                - /bin/sh
                - -c
                - date; echo Hello from the Kubernetes cluster
            restartPolicy: OnFailure
            kubectl apply -f date-job.yaml
            //Verify
```

```
kubectl get cj date-job -o yaml
B. CronJob Syntax:
* --> Minute
* --> Hours
* --> Day of The Month
* --> Month
* --> Day of the Week
*/1 * * * * --> Execute a command every one minutes.
vim date-job.yaml
apiVersion: batch/v1beta1
kind: CronJob
metadata:
name: date-job
spec:
schedule: "*/1 * * * *"
jobTemplate:
spec:
template:
- /bin/sh
--c
- date; echo Hello from the Kubernetes cluster
restartPolicy: OnFailure
kubectl apply -f date-job.yaml
//Verify
kubectl get cj date-job -o yaml
```

Answer: A

Question: 74

Create a daemonset named “Prometheus-monitoring” using image=prom/Prometheus which runs in all the nodes in the cluster. Verify the pod running in all the nodes

```
A. vim promo-ds.yaml
apiVersion: apps/v1
kind: DaemonSet
metadata:
name: prometheus-monitoring
spec:
selector:
matchLabels:
name: prometheus
template:
metadata:
labels:
name: prometheus
spec:
tolerations:
```

```
# remove it if your masters can't run pods
- key: node-role.kubernetes.io/master
effect: NoSchedule
containers:
- name: prometheus-container
image: prom/prometheus
volumeMounts:
- name: varlog
mountPath: /var/log
- name: varlibdockercontainers
mountPath: /var/lib/docker/containers
readOnly: true
volumes:
- name: varlog
emptyDir: {}
- name: varlibdockercontainers
emptyDir: {}
kubectl apply -f promo-ds.yaml
NOTE: Deamonset will get scheduled to "default" namespace, to
schedule deamonset in specific namespace, then add
"namespace" field in metadata
//Verify
kubectl get ds
NAME DESIRED CURRENT READY UP-TO-DATE
AVAILABLE NODE SELECTOR AGE
prometheus-monitoring 6 6 0 6
0 <none> 7s
kubectl get no # To get list of nodes in the cluster
// There are 6 nodes in the cluster, so a pod gets scheduled to
each node in the cluster
B. vim promo-ds.yaml
apiVersion: apps/v1
kind: DaemonSet
metadata:
name: prometheus-monitoring
spec:
selector:
matchLabels:
name: prometheus
template:
metadata:
labels:
name: prometheus
spec:
tolerations:
# remove it if your masters can't run pods
- key: node-role.kubernetes.io/master
effect: NoSchedule
```

```

containers:
- name: prometheus-container
- name: varlibdockercontainers
mountPath: /var/lib/docker/containers
readOnly: true
volumes:
- name: varlog
emptyDir: {}
- name: varlibdockercontainers
emptyDir: {}
kubectl apply -f promo-ds.yaml
NOTE: Deamonset will get scheduled to “default” namespace, to
schedule deamonset in specific namespace, then add
“namespace” field in metadata
//Verify
kubectl get ds
NAME DESIRED CURRENT READY UP-TO-DATE
AVAILABLE NODE SELECTOR AGE
prometheus-monitoring 8 8 0 6
0 <none> 7s
kubectl get no # To get list of nodes in the cluster
// There are 6 nodes in the cluster, so a pod gets scheduled to
each node in the cluster

```

Answer: A

Question: 75

CORRECT TEXT

Get list of persistent volumes and persistent volume claim in
the cluster

Answer:

kubectl get pv
kubectl get pvc

Question: 76

Create a redis pod, and have it use a non-persistent storage
(volume that lasts for the lifetime of the Pod)

A. vim redis-pod-vol.yaml

```

apiVersion: v1
kind: Pod
metadata:
name: redis
spec:

```

```
containers:  
- name: redis  
image: redis  
ports:  
- containerPort: 6379  
volumeMounts:  
- mountPath: /data  
name: redis-storage  
volumes:  
- name: redis-storage  
emptyDir: {}  
kubectl apply -f redis-pod-vol.yaml
```

```
B. vim redis-pod-vol.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
name: redis  
spec:  
containers:  
- name: redis  
image: redis  
ports:  
- containerPort: 6679  
volumeMounts:  
- mountPath: /data  
name: redis-storage  
volumes:  
- name: redis-storage  
emptyDir: {}  
kubectl apply -f redis-pod-vol.yaml
```

Answer: A

Question: 77

Create a hostPath PersistentVolume named task-pv-volume with storage 10Gi, access modes ReadWriteOnce, storageClassName manual, and volume at /mnt/data and verify

```
A. vim task-pv-volume.yaml  
apiVersion: v1  
kind: PersistentVolume  
metadata:  
name: task-pv-volume  
labels:  
type: local  
spec:  
storageClassName: ""
```

```
capacity:  
storage: 5Gi  
accessModes:  
- ReadWriteOnce  
hostPath:  
path: "/mnt/data"  
kubectl apply -f task-pv-volume.yaml  
//Verify  
kubectl get pv  
NAME CAPACITY ACCESS  
MODES RECLAIM POLICY STATUS CLAIM  
STORAGECLASS REASON AGE  
task-pv-volume 5Gi RWO  
Retain Available  
3s  
B. vim task-pv-volume.yaml  
apiVersion: v1  
kind: PersistentVolume  
metadata:  
name: task-pv-volume  
labels:  
type: local  
spec:  
storageClassName: ""  
capacity:  
storage: 5Gi  
accessModes:  
- ReadWriteOnce  
hostPath:  
path: "/mnt/data"  
kubectl apply -f task-pv-volume.yaml  
//Verify  
kubectl get pv  
NAME CAPACITY ACCESS  
MODES RECLAIM POLICY STATUS CLAIM  
STORAGECLASS REASON AGE  
task-pv-volume 4Gi RWO  
Retain Available  
8s
```

Answer: A

Question: 78

Create a PersistentVolumeClaim of at least 3Gi storage and access mode ReadWriteOnce and verify status is Bound

A. vim task-pv-claim.yaml

```
apiVersion: v2
kind: PersistentVolumeClaim
metadata:
  name: task-pv-claim
spec:
  storageClassName: ""
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 4Gi
  kubectl apply -f task-pv-claim.yaml
//Verify
  kubectl get pv
  NAME CAPACITY ACCESS
  MODES RECLAIM POLICY STATUS CLAIM
  STORAGECLASS REASON AGE
  task-pv-volume 4Gi RWO
  Retain Bound default/task-pv-claim
  6m16s
  kubectl get pvc
  NAME STATUS VOLUME
  CAPACITY ACCESS MODES STORAGECLASS AGE
  task-pv-claim Bound task-pv-volume
  5Gi RWO 6s
B. vim task-pv-claim.yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: task-pv-claim
spec:
  storageClassName: ""
  accessModes:
    - ReadWriteOnce
  resources:
    requests:
      storage: 3Gi
  kubectl apply -f task-pv-claim.yaml
//Verify
  kubectl get pv
  NAME CAPACITY ACCESS
  MODES RECLAIM POLICY STATUS CLAIM
  STORAGECLASS REASON AGE
  task-pv-volume 5Gi RWO
  Retain Bound default/task-pv-claim
  6m16s
  kubectl get pvc
  NAME STATUS VOLUME
```

CAPACITY ACCESS MODES STORAGECLASS AGE
task-pv-claim Bound task-pv-volume
5Gi RWO 6s

Answer: B

Question: 79

CORRECT TEXT

Delete persistent volume and persistent volume claim

Answer:

```
kubectl delete pvc task-pv-claim
kubectl delete pv task-pv-volume
// Verify
Kubectl get pv,pvc
```

Question: 80

Create a redis pod named "test-redis" and exec into that pod and create a file named "test-file.txt" with the text 'This is called the test file' in the path /data/redis and open another tab and exec again with the same pod and verifies file exist in the same path.

```
A. vim test-redis.yaml
apiVersion: v1
kind: Pod
metadata:
name: test-redis
spec:
containers:
- name: redis
image: redis
ports:
- containerPort: 6379
volumeMounts:
- mountPath: /data/redis
name: redis-storage
volumes:
- name: redis-storage
emptyDir: {}
kubectl apply -f redis-pod-vol.yaml
// first terminal
kubectl exec -it test-redis /bin/sh
cd /data/redis
echo 'This is called the test file' > file.txt
//open another tab
kubectl exec -it test-redis /bin/sh
```

```

cat /data/redis/file.txt
B. vim test-redis.yaml
apiVersion: v1
kind: Pod
metadata:
name: test-redis
spec:
containers:
- name: redis
image: redis
ports:
- containerPort: 6379
volumeMounts:
- mountPath: /data/redis
name: redis-storage
volumes:
kubectl exec -it test-redis /bin/sh
cd /data/redis
echo 'This is called the test file' > file.txt
//open another tab
kubectl exec -it test-redis /bin/sh
cat /data/redis/file.txt

```

Answer: A

Question: 81

CORRECT TEXT

Delete the above pod and create again from the same yaml file and verifies there is no "test-file.txt" in the path /data/redis
(Since non-persistent storage "emptyDir" is used).

Answer:

```

kubectl delete pod test-redis
kubectl create -f test-redis.yaml
kubectl exec -it test-redis /bin/sh
cat /data/redis/file.txt // file doesn't exist

```

Question: 82

Create PersistentVolume named task-pv-volume with storage 10Gi, access modes ReadWriteMany, storageClassName manual, and volume at /mnt/data and Create a PersistentVolumeClaim of at least 3Gi storage and access mode ReadWriteOnce and verify

```

A. vim task-pv-volume.yaml
apiVersion: v1
kind: PersistentVolume
metadata:

```

```
name: task-pv-volume
labels:
type: local
spec:
storageClassName: manual
capacity:
storage: 10Gi
accessModes:
- ReadWriteMany
hostPath:
path: "/mnt/data"
kubectl apply -f task-pv-volume.yaml
//Verify
kubectl get pv
vim task-pvc-volume.yaml
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
name: task-pv-claim
spec:
storageClassName: manual
accessModes:
- ReadWriteMany
resources:
requests:
storage: 3Gi
kubectl apply -f task-pvc-volume.yaml
//Verify
Kuk kubectl get pvc
B. vim task-pv-volume.yaml
apiVersion: v1
kind: PersistentVolume
metadata:
name: task-pv-volume
labels:
type: local
spec:
storageClassName: manual
capacity:
storage: 10Gi
accessModes:
- ReadWriteMany
hostPath:
path: "/mnt/data"
kubectl apply -f task-pv-volume.yaml
//Verify
kubectl get pv
vim task-pvc-volume.yaml
```

```
apiVersion: v1
- ReadWriteMany
resources:
requests:
storage: 3Gi
kubectl apply -f task-pvc-volume.yaml
//Verify
Kuk kubectl get pvc
```

Answer: A

Question: 83

Create an nginx pod with containerPort 80 and with a PersistentVolumeClaim "task-pv-claim" and has a mouth path "/usr/share/nginx/html"

A. vim nginx-pvc-pod.yaml

```
apiVersion: v1
kind: Pod
metadata:
name: task-pv-pod
spec:
volumes:
- name: task-pv-storage
persistentVolumeClaim:
claimName: task-pv-claim
containers:
- name: task-pv-container
image: nginx
ports:
- containerPort: 60
name: "http"
volumeMounts:
- mountPath: "/usr/share/nginx/html"
name: task-pv-storage
kubectl apply -f nginx-pvc-pod.yaml
// Verify
kubectl describe po "POD-Name" | grep -i volumes -A4
Volumes:
task-pv-storage:
Type: PersistentVolumeClaim (a reference to a
PersistentVolumeClaim in the same namespace)
ClaimName: task-pv-claim
ReadOnly: false
B. vim nginx-pvc-pod.yaml
```

apiVersion: v1
kind: Pod
metadata:

```

name: task-pv-pod
spec:
volumes:
- name: task-pv-storage
persistentVolumeClaim:
claimName: task-pv-claim
containers:
- name: task-pv-container
image: nginx
ports:
- containerPort: 80
name: "http"
volumeMounts:
- mountPath: "/usr/share/nginx/html"
name: task-pv-storage
kubectl apply -f nginx-pvc-pod.yaml
// Verify
kubectl describe po "POD-Name" | grep -i volumes -A5
Volumes:
task-pv-storage:
Type: PersistentVolumeClaim (a reference to a
PersistentVolumeClaim in the same namespace)
ClaimName: task-pv-claim
ReadOnly: false

```

Answer: B

Question: 84

CORRECT TEXT

Get list of PVs and order by size and write to file - /opt/pvlist.txt

kubectl get pv --sort-by=.spec.capacity.storage > /opt/pvlist.txt

Answer:

Question: 85

CORRECT TEXT

List all configmap and secrets in the cluster in all namespace and write it to a file /opt/configmap-secret

kubectl get configmap,secrets --all-namespaces >
/opt/configmap-secret
// Verify
Cat /opt/configmap-secret

Answer:

Question: 86

Create a configmap called myconfigmap with literal value
appname=myapp

A. kubectl create cm myconfigmap --from-literal=appname=myapp
// Verify
kubectl get cm -o yaml
(or)
kubectl describe cm

B. kubectl create cm myconfigmap --from-literal=appname=myapp
// Verify
(or)
kubectl describe cm

Answer: A

Question: 87

Create a file called “config.txt” with two values key1=value1
and key2=value2. Then create a configmap named “keyvalcfgmap” and read data from the file
“config.txt” and verify that configmap is
created correctly

A. cat >> config.txt << EOF
key1=value1
key2=value2
EOF
kubectl create cm keyvalcfgmap --from-file=config.txt
//Verify
kubectl get cm keyvalcfgmap -o yaml

B. cat >> config.txt << EOF
key1=value1
key2=value2
EOF
cat config.txt
// Create configmap from "config.txt" file
kubectl create cm keyvalcfgmap --from-file=config.txt
//Verify
kubectl get cm keyvalcfgmap -o yaml

Answer: B

Question: 88

Create an nginx pod and load environment values from the above configmap "keyvalcfgmap" and exec into the pod and verify the environment variables and delete the pod

A. // first run this command to save the pod yaml

```
kubectl run nginx --image=nginx --restart=Always --dry-run -o
```

```
yaml > nginx-pod.yaml
```

// edit the yml to below file and create

```
vim nginx-pod.yaml
```

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
labels:
```

```
run: nginx
```

```
name: nginx
```

```
spec:
```

```
containers:
```

```
- image: nginx
```

```
name: nginx
```

```
envFrom:
```

```
- configMapRef:
```

```
name: keyvalcfgmap
```

```
restartPolicy: Always
```

```
kubectl apply -f nginx-pod.yaml
```

// verify

```
kubectl exec -it nginx -- env
```

```
kubectl delete po nginx
```

B. // first run this command to save the pod yaml

```
kubectl run nginx --image=nginx --restart=Always --dry-run -o
```

```
yaml > nginx-pod.yaml
```

// edit the yml to below file and create

```
vim nginx-pod.yaml
```

```
apiVersion: v1
```

```
name: nginx
```

```
envFrom:
```

```
- configMapRef:
```

```
name: keyvalcfgmap
```

```
restartPolicy: Always
```

```
kubectl apply -f nginx-pod.yaml
```

// verify

```
kubectl exec -it nginx -- env
```

```
kubectl delete po nginx
```

Answer: A

Question: 89

Create a redis pod and mount “redis-config” as “redis.conf” inside redis container, name the config volume as “redis-volume”
redis-config path - /opt/redis-config

- A. 1
- B. 2
- C. Pending

Answer: C

Question: 90

```
// Create a configmap
kubectl create configmap redis-config --from-file=/opt/redisconfig
// Verify
kubectl get configmap redis-config -o yaml
// first run this command to save the pod yaml
kubectl run redis-pod --image=redis --restart=Always --dry-run
-o yaml > redis-pod.yaml
// edit the yaml to below file and create
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
    - name: redis
      image: redis
      env:
        - name: MASTER
          value: "true"
      ports:
        - containerPort: 6379
      volumeMounts:
        - mountPath: /redis-master-data
          name: data
        - mountPath: /redis-master
          name: config
      volumes:
        - name: data
        - name: configMap:
            name: example-redis-config
```

A. items:

```
- key: redis-config
path: redis.conf
cf
kk kubectl apply -f redis-pod.yml
// // Verify
K kubectl exec -it redis – cat /redis-master-data/redis.conf
```

B. items:

```
- key: redis-config
path: redis.conf
cf
// // Verify
K kubectl exec -it redis – cat /redis-master-data/redis.conf
```

Answer: A

Question: 91

Create a configmap called cfgvolume with values var1=val1, var2=val2 and create an nginx pod with volume nginx-volume which reads data from this configmap cfgvolume and put it on the path /etc/cfg

```
A. // first create a configmap cfgvolume
kubectl create cm cfgvolume --from-literal=var1=val1 --fromliteral=var2=val2
// verify the configmap
kubectl describe cm cfgvolume
// create the config map
kubectl create -f nginx-volume.yml
vim nginx-configmap-pod.yaml
apiVersion: v1
kind: Pod
- name: nginx-volume
configMap:
name: cfgvolume
containers:
- image: nginx
name: nginx
volumeMounts:
- name: nginx-volume
mountPath: /etc/cfg
restartPolicy: Always
k kubectl apply -f nginx-configmap-pod.yaml
// // Verify
// exec into the pod
kubectl exec -it nginx -- /bin/sh
// check the path
```

```

cd /etc/cfg
B. // first create a configmap cfgvolume
kubectl create cm cfgvolume --from-literal=var1=val1 --fromliteral=var2=val2
// verify the configmap
kubectl describe cm cfgvolume
// create the config map
kubectl create -f nginx-volume.yml
vim nginx-configmap-pod.yaml
apiVersion: v1
kind: Pod
metadata:
labels:
run: nginx
name: nginx
spec:
volumes:
- name: nginx-volume
configMap:
name: cfgvolume
containers:
- image: nginx
name: nginx
volumeMounts:
- name: nginx-volume
mountPath: /etc/cfg
restartPolicy: Always
k kubectl apply -f nginx-configmap-pod.yaml
// Verify
// exec into the pod
kubectl exec -it nginx -- /bin/sh
// check the path
cd /etc/cfg

```

Answer: B

Question: 92

Create a secret mysecret with values user=myuser and password=mypassword

```

A. kubectl create secret generic my-secret --fromliteral=username=user --fromliteral=password=mypassword
// Verify
kubectl get secret --all-namespaces
kubectl get secret generic my-secret -o yaml
B. kubectl create secret generic my-secret --fromliteral=username=user --fromliteral=password=mypassword
// Verify
kubectl get secret generic my-secret -o yaml

```

Answer: A

Question: 93

Create an nginx pod which reads username as the environment variable

A. // create a yml file
kubectl run nginx --image=nginx --restart=Never --dry-run -o
yaml > nginx.yml
// add env section below and create
apiVersion: v1
kind: Pod
metadata:
labels:
run: nginx
name: nginx
spec:
containers:
- image: nginx
name: nginx
env:
- name: USER_NAME
restartPolicy: Never
kubectl create -f nginx.yml
//Verify
kubectl exec -it nginx – env
B. // create a yml file
kubectl run nginx --image=nginx --restart=Never --dry-run -o
yaml > nginx.yml
// add env section below and create
apiVersion: v1
kind: Pod
metadata:
labels:
run: nginx
name: nginx
spec:
containers:
- image: nginx
name: nginx
env:
- name: USER_NAME
valueFrom:
secretKeyRef:
name: my-secret
key: username
restartPolicy: Never

```
kubectl create -f nginx.yml  
//Verify  
kubectl exec -it nginx – env
```

Answer: B

Question: 94

Create an nginx pod which loads the secret as environment variables

A. // create a yaml file

```
kubectl run nginx --image=nginx --restart=Never --dry-run -o  
yaml > nginx.yml  
// add env section below and create  
vim nginx.yaml  
apiVersion: v1  
kind: Pod  
metadata:  
labels:  
run: nginx  
name: nginx  
spec:  
containers:  
- image: nginx  
name: nginx  
envFrom:  
- secretRef:  
name: my-secret  
restartPolicy: Never  
kubectl apply -f nginx.yaml  
//verify  
kubectl exec -it nginx – env
```

B. // create a yaml file

```
kubectl run nginx --image=nginx --restart=Never --dry-run -o  
yaml > nginx.yml  
// add env section below and create  
vim nginx.yaml  
run: nginx  
name: nginx  
spec:  
containers:  
- image: nginx  
name: nginx  
envFrom:  
- secretRef:  
name: my-secret  
restartPolicy: Never  
kubectl apply -f nginx.yaml
```

```
//verify  
kubectl exec -it nginx - env
```

Answer: A

Question: 95

List all service account and create a service account called “admin”

A. kubectl get sa
kubectl get sa --all-namespaces
kubectl create sa admin
//Verify
kubectl get sa admin -o yaml
B. kubectl get sa
kubectl get sa --all-namespaces
//Verify
kubectl get sa admin -o yaml

Answer: A

Question: 96

Create a busybox pod which executes this command sleep 3600 with the service account admin and verify

A. kubectl run busybox --image=busybox --restart=Always --dry-run
-o yaml -- /bin/sh -c "sleep 3600" > busybox.yml
// Edit busybox.yaml file
apiVersion: v1
kind: Pod
metadata:
creationTimestamp: null
labels:
run: busybox
name: busybox
spec:
serviceAccountName: admin
containers:
- args:
- /bin/sh
- -c
- sleep 3600
image: busybox
name: busybox
restartPolicy: Always
// verify

```
K kubectl describe po busybox
B. kubectl run busybox --image=busybox --restart=Always --dry-run
-o yaml -- /bin/sh -c "sleep 3600" > busybox.yaml
// Edit busybox.yaml file
apiVersion: v1
kind: Pod
metadata:
creationTimestamp: null
labels:
run: busybox
name: busybox
spec:
serviceAccountName: admin
containers:
- args:
- /bin/sh
- -c
- sleep 3800
image: busybox
name: busybox
restartPolicy: Always
// verify
K kubectl describe po busybox
```

Answer: A

Question: 97

Create an nginx pod with container Port 80 and it should only receive traffic only if it checks the endpoint / on port 80 and verify and delete the pod.

```
A. kubectl run nginx --image=nginx --restart=Never --port=80 --
dry-run -o yaml > nginx-pod.yaml
// add the readinessProbe section and create
vim nginx-pod.yaml
run: nginx
name: nginx
spec:
containers:
- image: nginx
name: nginx
ports:
- containerPort: 60
readinessProbe:
httpGet:
path: /
port: 60
restartPolicy: Never
```

```
kubectl apply -f nginx-pod.yaml
// verify
kubectl describe pod nginx | grep -i readiness
kubectl delete po nginx
B. kubectl run nginx --image=nginx --restart=Never --port=80 --
dry-run -o yaml > nginx-pod.yaml
// add the readinessProbe section and create
vim nginx-pod.yaml
apiVersion: v1
kind: Pod
metadata:
labels:
run: nginx
name: nginx
spec:
containers:
- image: nginx
name: nginx
ports:
- containerPort: 80
readinessProbe:
httpGet:
path: /
port: 80
restartPolicy: Never
kubectl apply -f nginx-pod.yaml
// verify
kubectl describe pod nginx | grep -i readiness
kubectl delete po nginx
```

Answer: B

Question: 98

List all the events sorted by timestamp and put them into file.log and verify

- A. kubectl get events --sort-by=.metadata.creationTimestamp
kubectl get events --sort-by=.metadata.creationTimestamp >
test-file.log
cat test-file.log
- B. kubectl get events --sort-by=.metadata.creationTimestamp
// putting them into file.log
kubectl get events --sort-by=.metadata.creationTimestamp >
cat test-file.log
- C. kubectl get events --sort-by=.metadata.creationTimestamp
// putting them into file.log
kubectl get events --sort-by=.metadata.creationTimestamp >
test-file.log

```
cat test-file.log
```

Answer: C

Question: 99

Get the memory and CPU usage of all the pods and find out top 3 pods which have the highest usage and put them into the cpuusage.txt file

```
A. // Get the top 3 pods
kubectl top pod --all-namespaces | sort --reverse --key 3 --
numeric | head -3
// putting into file
kubectl top pod --all-namespaces | sort --reverse --key 3 --
numeric | head -3 > cpu-usage.txt
// verify
cat cpu-usage.txt
B. // Get the top 3 pods
kubectl top pod --all-namespaces | sort --reverse --key 3 --
numeric | head -8
// putting into file
kubectl top pod --all-namespaces | sort --reverse --key 6 --
numeric | head -6 > cpu-usage.txt
// verify
cat cpu-usage.txt
```

Answer: A

Question: 100

CORRECT TEXT

Create the service as type NodePort with the port 32767 for the nginx pod with the pod selector app: my-nginx

Answer:

```
kubectl run nginx --image=nginx --restart=Never --
labels=app=nginx --port=80 --dry-run -o yaml > nginx-pod.yaml
```

Question: 101

Create a Pod nginx and specify a CPU request and a CPU limit of 0.5 and 1 respectively.

```
A. // create a yml file
kubectl run nginx-pod --image=nginx --restart=Never --dry-run -
o yaml > nginx-pod.yaml
// add the resources section and create
vim nginx-pod.yaml
```

```
apiVersion: v1
kind: Pod
metadata:
labels:
run: nginx
name: nginx
spec:
containers:
- image: nginx
name: nginx
resources:
requests:
cpu: "0.5"
limits:
cpu: "1"
restartPolicy: Always
kubectl apply -f nginx-pod.yaml
// verify
kubectl top pod
B. // create a yml file
kubectl run nginx-pod --image=nginx --restart=Never --dry-run -
o yaml > nginx-pod.yml
// add the resources section and create
vim nginx-pod.yaml
apiVersion: v1
kind: Pod
metadata:
labels:
run: nginx
name: nginx
spec:
containers:
- image: nginx
name: nginx
resources:
requests:
cpu: "0.4"
limits:
cpu: "1"
restartPolicy: Always
kubectl apply -f nginx-pod.yaml
// verify
kubectl top pod
```

Answer: A

Question: 102

Create a Pod nginx and specify both CPU, memory requests and limits together and verify.

```
A. kubectl run nginx-request --image=nginx --restart=Always --dryrun -o yaml > nginx-request.yaml
// add the resources section and create
apiVersion: v1
kind: Pod
metadata:
labels:
run: nginx
name: nginx
resources:
requests:
memory: "100Mi"
cpu: "0.4"
limits:
memory: "200Mi"
cpu: "7"
restartPolicy: Always
k kubectl apply -f nginx-request.yaml
// Verify
Kubectl top po
B. kubectl run nginx-request --image=nginx --restart=Always --dryrun -o yaml > nginx-request.yaml
// add the resources section and create
apiVersion: v1
kind: Pod
metadata:
labels:
run: nginx
name: nginx-request
spec:
containers:
- image: nginx
name: nginx
resources:
requests:
memory: "100Mi"
cpu: "0.5"
limits:
memory: "200Mi"
cpu: "1"
restartPolicy: Always
k kubectl apply -f nginx-request.yaml
// Verify
Kubectl top po
```

Answer: B

Question: 103

Set CPU and memory requests and limits for existing pod name "nginx-prod".

Set requests for CPU and Memory as 100m and 256Mi respectively

Set limits for CPU and Memory as 200m and 512Mi respectively

A. kubectl get po

```
kubectl set resources po nginx-prod --
```

```
limits(cpu=200m, memory=512Mi) --requests(cpu=100m, memory=256Mi)
```

//Verify

```
kubectl top po
```

```
kubectl describe po nginx-prod
```

B. kubectl get po

```
kubectl set resources po nginx-prod --
```

```
limits(cpu=200m, memory=512Mi) --requests(cpu=100m, memory=256Mi)
```

//Verify

```
kubectl describe po nginx-prod
```

Answer: A

Question: 104

Create an nginx pod with containerPort 80 and it should check the pod running at endpoint / healthz on port 80 and verify and delete the pod.

A. kubectl run nginx --image=nginx --restart=Always --port=80 --

```
dry-run -o yaml > nginx-pod.yaml
```

// add the livenessProbe section and create

```
apiVersion: v1
```

```
kind: Pod
```

```
metadata:
```

```
labels:
```

```
run: nginx
```

```
name: nginx
```

```
spec:
```

```
containers:
```

```
- image: nginx
```

```
name: nginx
```

```
ports:
```

```
- containerPort: 80
```

```
livenessProbe:
```

```
httpGet:
```

```
path: /healthz
```

```
port: 80
```

```
restartPolicy: Always
```

```
kubectl create -f nginx-pod.yaml
```

// verify

```
kubectl describe pod nginx | grep -i readiness
```

```
kubectl delete po nginx
B. kubectl run nginx --image=nginx --restart=Always --port=80 --
dry-run -o yaml > nginx-pod.yaml
// add the livenessProbe section and create
apiVersion: v1
kind: Pod
metadata:
labels:
containers:
- image: nginx
name: nginx
ports:
- containerPort: 60
livenessProbe:
httpGet:
path: /healthz
port: 60
restartPolicy: Always
kubectl create -f nginx-pod.yaml
// verify
kubectl describe pod nginx | grep -i readiness
kubectl delete po nginx
```

Answer: A

Question: 105

Create a NetworkPolicy which denies all ingress traffic

```
A. apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
name: default-deny
spec:
podSelector: {}
policyTypes:
- Ingress
```

```
B. apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
name: default-deny
spec:
podSelector: {}
policyTypes:
- Ingress
```

Answer: A

Question: 106

Allow traffic from all the pods in "web" namespace and from pods with label "type=monitoring" to the pods matching label "app: db"

A. kubectl create namespace web

```
kubectl label namespace/web app=web
vim web-allow-all-ns-monitoring.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
name: web-allow-all-ns-monitoring
namespace: default
spec:
podSelector:
matchLabels:
app: db
ingress:
- from:
- namespaceSelector:
matchLabels:
app: web
podSelector:
matchLabels:
type: monitoring
k kubectl apply -f web-allow-all-ns-monitoring.yaml
```

B. kubectl create namespace web

```
kubectl label namespace/web app=web
vim web-allow-all-ns-monitoring.yaml
apiVersion: networking.k8s.io/v1
kind: NetworkPolicy
metadata:
name: web-allow-all-ns-monitoring
namespace: default
spec:
podSelector:
podSelector:
matchLabels:
type: monitoring
k kubectl apply -f web-allow-all-ns-monitoring.yaml
```

Answer: A

Question: 107

Create a ETCD backup of kubernetes cluster

Note : You don't need to memorize command, refer -

<https://kubernetes.io/docs/tasks/administer-cluster/configureupgrade-etcd/> during exam

A. ETCDCTL_API=3 etcdctl --endpoints=[ENDPOINT] --cacert=[CA CERT]
--cert=[ETCD SERVER CERT] --key=[ETCD SERVER KEY] snapshot save
[BACKUP FILE NAME]

In exam, cluster setup is done with kubeadm , this means ETCD used by the kubernetes cluster is coming from static pod.

kubectl get pod -n kube-system

kubectl describe pod etcd-master -n kube-system

You can locate the information on

endpoint: — advertise-client-urls=https://172.17.0.15:2379

ca certificate: — trusted-cafile=/etc/kubernetes/pki/etcd/ca.crt

server certificate : — certfile=/etc/kubernetes/pki/etcd/server.crt

key: — key-file=/etc/kubernetes/pki/etcd/server.key

To Create backup

export ETCDCTL_API=3

(or)

```
ETCDCTL_API=3 etcdctl ETCDCTL_API=3 etcdctl --
endpoints=https://172.17.0.15:2379 --
cacert=/etc/kubernetes/pki/etcd/ca.crt --
cert=/etc/kubernetes/pki/etcd/server.crt --
key=/etc/kubernetes/pki/etcd/server.key snapshot save etcdsnapshot.db
//Verify
```

```
ETCDCTL_API=3 etcdctl --write-out=table snapshot status
snapshot.db
```

B. ETCDCTL_API=3 etcdctl --endpoints=[ENDPOINT] --cacert=[CA CERT]
--cert=[ETCD SERVER CERT] --key=[ETCD SERVER KEY] snapshot save
[BACKUP FILE NAME]

In exam, cluster setup is done with kubeadm , this means ETCD used by the kubernetes cluster is coming from static pod.

kubectl get pod -n kube-system

kubectl describe pod etcd-master -n kube-system

You can locate the information on

endpoint: — advertise-client-urls=https://172.16.0.18:2379

ca certificate: — trusted-cafile=/etc/kubernetes/pki/etcd/ca.crt

server certificate : — certfile=/etc/kubernetes/pki/etcd/server.crt

key: — key-file=/etc/kubernetes/pki/etcd/server.key

To Create backup

export ETCDCTL_API=3

(or)

```
ETCDCTL_API=3 etcdctl ETCDCTL_API=3 etcdctl --
endpoints=https://172.17.0.15:2379 --
key=/etc/kubernetes/pki/etcd/server.key snapshot save etcdsnapshot.db
//Verify
```

```
ETCDCTL_API=3 etcdctl --write-out=table snapshot status
```

snapshot.db

Answer: A

Question: 108

create a pod in a specific node (node1) by placing the pod definition file in a particular folder “/etc/kubernetes/manifests”.

A. Generate YAML before we SSH to the specific node, then copy the YAML into the exam notepad to use it after SSH into worker node.

SSH to the node: "ssh node1"

Gain admin privileges to the node: "sudo -i"

Move to the manifest-path "cd /etc/kubernetes/manifests"

kubelet config file -- /var/lib/kubelet/config.yaml

Edit the config file “vi /var/lib/kubelet/config.yaml” to add staticPodPath

staticPodPath: /etc/kubernetes/manifests

Restart the kubelet “systemctl restart kubelet”

B. Generate YAML before we SSH to the specific node, then copy the YAML into the exam notepad to use it after SSH into worker node.

SSH to the node: "ssh node1"

Gain admin privileges to the node: "sudo -i"

Move to the manifest-path "cd /etc/kubernetes/manifests"

Place the generated YAML into the folder "vi nginx.yaml"

Find the kubelet config file path "ps -aux | grep kubelet" . This will output information on kubelet process. Locate the kubelet config file location.

kubelet config file -- /var/lib/kubelet/config.yaml

Edit the config file “vi /var/lib/kubelet/config.yaml” to add staticPodPath

staticPodPath: /etc/kubernetes/manifests

Restart the kubelet “systemctl restart kubelet”

Answer: B

Question: 109

CORRECT TEXT

View certificate details in /etc/kubernetes/pki

Answer:

// Verify Public Key / Certificate file

Openssl x509 -in certificate.crt -noout -text

// Verify Private Key

Openssl rsa -in “private-key” -check

// Verify CSR

`Openssl req -text -noout -verify`

Question: 110

CORRECT TEXT

Verify certificate expiry date for ca certificate in /etc/kubernetes/pki

Answer:

`openssl x509 -in ca.crt -noout -text | grep -i validity -A 4`

Question: 111

Evict all existing pods from a node-1 and make the node unschedulable for new pods.

A. `kubectl get nodes`

`kubectl drain node-1 #It will evict pods running on node-1 to other nodes in the cluster`

`kubectl cordon node-1 # New pods cannot be scheduled to the node`

// Verify

`kubectl get no`

When you cordon a node, the status shows “SchedulingDisabled”

B. `kubectl get nodes`

`kubectl drain node-1 #It will evict pods running on node-1 to other nodes in the cluster`

// Verify

`kubectl get no`

When you cordon a node, the status shows “SchedulingDisabled”

Answer: A

Question: 112

CORRECT TEXT

Make the node schedulable by uncordon the node

Answer:

`kubectl uncordon node-1`

//verify

`kubectl get no`

Question: 113

Check nodes which are ready and print it to a file /opt/nodestatus

A. `JSONPATH='{range .items[*]}{@.metadata.name}:{range @.status.conditions[*]}{@.type}={@.status};{end}{end}' \`
`&& kubectl get nodes -o jsonpath="$JSONPATH" | grep "Ready=True" > /opt/node-status`
`//Verify`

```
cat /opt/node-status
B. JSONPATH='{range .items[*]}{@.metadata.name}:{range
@.status.conditions[*]}{@.type}={@.status};{end}{end}' \
//Verify
cat /opt/node-status
```

Answer: A

Question: 114

Install a kubernetes cluster with one master and one worker using kubeadm

A. This is a straightforward question, you need to install kubernetes cluster using kubeadm with one master and one worker.

Refer : <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/>

B. This is a straightforward question, you need to install kubernetes cluster using kubeadm with one master and one worker.

Installation is considered success once both master and worker nodes become available.

Refer : <https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/>

Answer: B

Question: 115

CORRECT TEXT

Get list of PVs and order by size and write to file "/opt/pvstorage.txt"

Answer:

```
kubectl get pv --sort-by=.spec.capacity.storage > /opt/pv
storage.txt
```

Question: 116

Get the number of schedulable nodes and write to a file
/opt/schedulable-nodes.txt

```
A. kubectl get nodes -o jsonpath="{range
.items[*]}.{.metadata.name}
.spec.taints[?(@.effect=='NoSchedule')].effect}{\"\\n\"}{end}"
| awk 'NF==1 {print $0}' > /opt/schedulable-nodes.txt
// Verify
cat /opt/schedulable-nodes.txt
B. kubectl get nodes -o jsonpath="{range
.items[*]}.{.metadata.name}
.spec.taints[?(@.effect=='NoSchedule')].effect}{\"\\n\"}{end}"
| awk 'NF==11 {print $0}' > /opt/schedulable-nodes.txt
```

```
// Verify
cat /opt/schedulable-nodes.txt
```

Answer: A

Question: 117

Fix a node that shows as non-ready

A. Kubectl get nodes

```
// Check which node shows a not ready
kubectl describe nodes "node-name"
// Login to the node which shows as not ready and check the
process for kubelet, docker , kube-proxy.
// systemctl status kubelet (or) ps -aux | grep -i "processname"
// If the process is not started, then start using
systemctl start kubelet / docker
// Verify
ps -auxxww | grep -i "process-name"
kubectl get nodes
```

B. Kubectl get nodes

```
// Check which node shows a not ready
kubectl describe nodes "node-name"
// Login to the node which shows as not ready and check the
systemctl start kubelet / docker
// Verify
ps -auxxww | grep -i "process-name"
kubectl get nodes
```

Answer: A

Question: 118

CORRECT TEXT

Label a node as app=test and verify

```
kubectl label node node-name app=test
// Verify
kubectl get no --show-labels
kubectl get no -l app=test
```

Answer:

Question: 119

Deploy a pod with image=redis on a node with label disktype:ssd

A. // Get list of nodes

```

kubectl get nodes
//Get node with the label disktype=ssd
kubectl get no -l disktype=ssd
// Create a sample yaml file
kubectl run node-redis --generator=run-pod/v1 --image=redis --dry
run -o yaml > test-redis.yaml
// Edit test-redis.yaml file and add nodeSelector
vim test-redis.yaml
apiVersion: v1
kind: Pod
metadata:
name: redis
spec:
nodeSelector:
disktype: ssd
containers:
- name: node-redis
image: redis
imagePullPolicy: IfNotPresent
kubectl apply -f test-redis.yaml
/// Verify
K kubectl get po -o wide

```

B. // Get list of nodes

```

kubectl get nodes
//Get node with the label disktype=ssd
kubectl get no -l disktype=ssd
// Create a sample yaml file
kubectl run node-redis --generator=run-pod/v1 --image=redis --dry
run -o yaml > test-redis.yaml
// Edit test-redis.yaml file and add nodeSelector
vim test-redis.yaml
apiVersion: v1
- name: node-redis
image: redis
imagePullPolicy: IfNotPresent
kubectl apply -f test-redis.yaml
/// Verify
K kubectl get po -o wide

```

Answer: A

Question: 120

Add a taint to node "worker-2" with effect as "NoSchedule" and list the node with taint effect as "NoSchedule"

A. // Add taint to node "worker-2"

```

kubectl taint nodes worker-2 key=value:NoSchedule
.items[*]{.metadata.name} {.spec.taints[?(@.effect=='NoSchedule' )].effect}{\"\\n\"}{end}" | awk 'NF==2{print $0}'
B. // Add taint to node "worker-2"
kubectl taint nodes worker-2 key=value:NoSchedule
// Verify
// Using "custom-columns" , you can customize which column to
be printed
kubectl get nodes -o customcolumns=NAME:.metadata.name,TAINTS:.spec.taints --no-headers
// Using jsonpath
kubectl get nodes -o jsonpath="{range .items[*]{.metadata.name} {.spec.taints[?(@.effect=='NoSchedule' )].effect}{\"\\n\"}{end}" | awk 'NF==2{print $0}'"

```

Answer: B

Question: 121

CORRECT TEXT

Remove taint added to node "worker-2"

Answer:

```

kubectl taint nodes worker-2 key:NoSchedule-
// Verify
You will see a message "node/worker-2 untainted"
kubectl get nodes -o customcolumns=NAME:.metadata.name,TAINTS:.spec.taints --no-headers

```

Question: 122

CORRECT TEXT

Print all pod name and all image name and write it to a file
name "/opt/pod-details.txt"

Answer:

```

kubectl get pods -o=custom-columns='Pod
Name:metadata.name','Image:spec.containers[*].image' >
/opt/pod-details.txt

```


Thank You for Purchasing CKA PDF

Bonus Products

3 SETS OF REAL EXAM SCREENSHOTS - 2020

Starting from next page , there are 3 complete sets of real exam screenshots taken from testing center. It is a bonus product alongwith the PDF With questions and answers that you have read above this page. All you have to do is to find the correct answers to these real exam screenshots before attempting the real exam.

First SET starts below

BEFORE YOU BEGIN: IMPORTANT INSTRUCTIONS:

Each question on this exam must be completed on a designated *cluster/configuration context*. To minimize switching, the questions are grouped so that all questions on a given *cluster* appear consecutively.

There are six *clusters* that comprise the exam environment, made up of varying numbers of *containers*, as follows:

- **k8s** - 1 master, 2 worker
- **hk8s** - 1 master, 2 worker
- **bk8s** - 1 master, 1 worker
- **wk8s** - 1 master, 2 worker
- **ek8s** - 1 master, 2 worker
- **ik8s** - 1 master, 1 base node

At the start of each question you'll be provided with the command to ensure you are on the correct *cluster* e.g.

Set configuration context: `$ kubectl config use-context k8s`

nodes making up each *cluster* can be reached via `ssh`, using a command such as the following:

varying numbers of containers, as follows.

- **k8s** - 1 master, 2 worker
- **hk8s** - 1 master, 2 worker
- **bk8s** - 1 master, 1 worker
- **wk8s** - 1 master, 2 worker
- **ek8s** - 1 master, 2 worker
- **ik8s** - 1 master, 1 base node

At the start of each question you'll be provided with the command to ensure you are on the correct cluster e.g.

Set configuration context: `$ kubectl config use-context k8s`

nodes making up each cluster can be reached via `ssh`, using a command such as the following:

```
$ ssh k8s-node-0
```

Elevated privileges can be assumed on any node with the following command:

```
$ sudo -i
```

When you have finished working on a node, you should return to the base node (with hostname `node-1`) before attempting any further

```
$ ssh k8s-node-0
```

Elevated privileges can be assumed on any *node* with the following command:

```
$ sudo -i
```

When you have finished working on a *node*, you should return to the *base node* (with hostname `node-1`) before attempting any further questions. Nested- `ssh` is not supported.

You can use `kubectl` and the appropriate context to work on any *cluster* from the *base node*. When connected to a *cluster Member* via `ssh`, you will only be able to work on that particular *cluster* via `kubectl`.

Further instructions for connecting to *cluster nodes* will be provided in the appropriate questions.

Please use the navigation panel to proceed to your first question.

Question: 0/24



Question 0→

Navigate All Questions



Question weight: 5%

Set configuration context: `$ kubectl config use-context k8s`

Monitor the logs of pod **foo** and:

- Extract log lines corresponding to error `unable-to-access-website`
- Write them to `/opt/KULM00201/foo`

Question: 1/24



Question 1 ▾

Navigate All Questions

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

List all persistent volumes sorted by **name**, saving the full `kubectl` output to `/opt/KUCC00102/pv_list`. Use `kubectl`'s own functionality for sorting the output, and do not manipulate it any further.

Question: 2/24



Question 2▼

[Navigate All Questions](#)



Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

Ensure a single instance of pod `nginx` is running on each *node* of the *Kubernetes cluster* where `nginx` also represents the *image name* which has to be used. Do not override any *taints* currently in place.

Use **DaemonSet** to complete this task and use `ds-kusc00201` as *DaemonSet name*.

Question: 3/24



Question 3▼

[Navigate All Questions](#)

Question weight: 7%

Set configuration context: `$ kubectl config use-context k8s`

Perform the following tasks:

- Add an *init container* to `lumpy-koala` (which has been defined in spec file `/opt/KUCC00108/pod-spec-KUCC00108.yaml`)
- The *init container* should create an empty file named `/workdir/eager.txt`
- If `/workdir/eager.txt` is not detected, the pod should exit
- Once the spec file has been updated with the *init container* definition, the pod should be created

Question: 4/24



Question 4▼

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context k8s`

Create a pod named `kucc4` with a single app container for each of the following *images* running inside (there may be between 1 and 4 *images* specified): `nginx + redis + memcached`.

Question: 5/24



Question 5▼

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context k8s`

Create a pod named `kucc4` with a single app container for each of the following *images* running inside (there may be between 1 and 4 *images* specified): `nginx + redis + memcached`.

Question: 5/24



Question 5▼

[Navigate All Questions](#)

Question weight: 2%

Set configuration context: `$ kubectl config use-context k8s`

Schedule a pod as follows:

- Name: `nginx-kusc00101`
- Image: `nginx`
- Node selector: `disk=spinning`

Question: 6/24



Question 6

[Navigate All Questions](#)

Question weight  4%

Set configuration context: `$ kubectl config use-context k8s`

Create a deployment as follows:

- Name: `nginx-app`
- Using container `nginx` with version `1.10.2-alpine`
- The deployment should contain `3` replicas

Next, deploy the application with new version `1.13.0-alpine`, by performing a rolling update, and record that update.

Finally, rollback that update to the previous version `1.10.2-alpine`.

Question: 7/24



Question 7 ▾

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context k8s`

Create and configure the service `front-end-service` so it's accessible through `ClusterIP` and routes to the existing `pod` named `front-end`.

Question: 8/24



Question 8▼

[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

Create a pod as follows:

- Name: `nginx`
- Using image: `nginx`
- In a new Kubernetes namespace named: `website-frontend`

Question: 9/24



Question 9▼

[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

Create a deployment spec file that will:

- Launch `7` replicas of the `redis` image with the label `app_runtime_stage=test`
- deployment name: `kual00201`

Save a copy of this spec file to `/opt/KUAL00201/spec_deploy.yaml` (or `.json`).

When you are done, clean up (delete) any new Kubernetes API object that you produced during this task.

Question: 10/24



Question 10▼

[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

Create a file: `/opt/KUCC00302/kucc00302.txt` that lists all pods that implement service **foo** in namespace **production**.

The format of the file should be one pod name per line.

Question: 11/24



Question 11

[Navigate All Questions](#)

Question weight: 9%

Set configuration context: `$ kubectl config use-context k8s`

Create a Kubernetes secret as follows:

- Name: `supersecret`
- username: `alice`

Create a pod named `pod-secrets-via-file`, using the `redis` image, which mounts a secret named `super-secret` at `/secrets`.

Create a second pod named `pod-secrets-via-env`, using the `redis` image, which exports `username` as `TOPSECRET`.

Question: 12/24



Question 12

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context k8s`

Create a pod as follows:

- Name: `non-persistent-redis`
- container image: `redis`
- Persistent volume with name: `app-cache`
- Mount path: `/data/redis`

It should launch in the `staging` namespace and the volume **must not** be persistent.

Question: 13/24



Question 13

[Navigate All Questions](#)

Question weight: 1%

Set configuration context: `$ kubectl config use-context k8s`

Scale the deployment **guestbook** to 3 pods.

Question: 14/24



Question 14▼

[Navigate All Questions](#)



Question weight: 2%

Set configuration context: `$ kubectl config use-context k8s`

Check to see how many nodes are ready (not including nodes tainted `NoSchedule`) and write the number to `/opt/KUCC00104/kucc00104.txt`.

Question: 15/24



Question 15▼

[Navigate All Questions](#)

Question weight: 2%

Set configuration context: `$ kubectl config use-context k8s`

From the *pod* /label **name=cpu-utilizer**, find *pods* running high CPU workloads and write the name of the *pod* consuming most CPU to the file `/opt/KUTR00102/KUTR00102.txt` (which already exists).

Question: 16/24



Question 16▼

[Navigate All Questions](#)

Question weight: 7%

Set configuration context: `$ kubectl config use-context k8s`

Create a deployment as follows:

- Name: `nginx-test`
- Exposed via a service `nginx-test`
- Ensure that the *service & pod* are accessible via their respective *DNS* records
- The *container(s)* within any *pod(s)* running as a part of this *deployment* should use the `nginx` image

Next, use the utility `nslookup` to look up the *DNS* records of the *service & pod* and write the output to `/opt/KUNW00601/service.dns` and `/opt/KUNW00601/pod.dns` respectively.

Question: 17/24



Question 17▼

[Navigate All Questions](#)

Question weight: 7%

No configuration context change required for this item

Create a snapshot of the `etcd` instance running at

`https://127.0.0.1:2379`, saving the snapshot to the file path
`/var/lib/backup/etcd-snapshot.db`.

The `etcd` instance is running `etcd` version 3.3.10.



The following *TLS* certificates/key are supplied for connecting to the server with `etcdctl`:

- CA certificate: `/opt/KUCM00302/ca.crt`
- Client certificate: `/opt/KUCM00302/etcd-client.crt`
- Client key: `/opt/KUCM00302/etcd-client.key`

Question: 18/24



Question 18

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context ek8s`

Set the *node* named `ek8s-node-0` as unavailable and reschedule all the *pods* running on it.

Question: 19/24



I

Question 19▼

Navigate All Questions

Question weight: 4%

Set configuration context: `$ kubectl config use-context wk8s`

A Kubernetes worker node, named **wk8s-node-0** is in state **NotReady**. Investigate why this is the case, and perform any appropriate steps to bring the *node* to a **Ready** state, ensuring that any changes are made permanent.

Hints:



- You can `ssh` to the failed node using: `$ ssh wk8s-node-0`
- You can assume elevated privileges on the *node* with the following command: `$ sudo -i`

Question: 20/24



Question 20

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context wk8s`

Configure the `kubelet` `systemd`-managed service, on the `node` labelled with `name=wk8s-node-1`, to launch a pod containing a single `container` of `image` `httpd` named `webtool` automatically. Any `spec` files required should be placed in the `/etc/kubernetes/manifests` directory on the `node`.

Hints:

- You can `ssh` to the appropriate `node` using: `$ ssh wk8s-node-1`
- You can assume elevated privileges on the `node` with the following command: `$ sudo -i`

Question: 21/24



Question 21▼

[Navigate All Questions](#)

Question weight: 8%

No configuration context change required for this item

Important: For this item, you will have to `ssh` to the nodes `ik8s-master-0` and `ik8s-node-0` and complete all tasks on these *nodes*.

Ensure that you return to the *base node* (hostname: `node-1`) when you have completed this item.

As an administrator of a small development team, you have been asked to set up a *Kubernetes cluster* to test the viability of a new *application*.

You must use `kubeadm` to perform this task. Any `kubeadm` invocations will require the use of the `--ignore-preflight-errors=all` option.

- Configure the *node* `ik8s-master-0` as a *master node*.
- Join the *node* `ik8s-node-0` to the *cluster*.

Docker is already installed on both *nodes* and `apt` has been configured so that you can install the required tools.

You must use the `kubeadm` configuration file located at
`/etc/kubeadm.conf` when initializing your *cluster*.

Ensure that you return to the *base node* (hostname: `node-1`) when you have completed this item.

As an administrator of a small development team, you have been asked to set up a *Kubernetes cluster* to test the viability of a new *application*.

You must use `kubeadm` to perform this task. Any `kubeadm` invocations will require the use of the `--ignore-preflight-errors=all` option.

- Configure the *node* `ik8s-master-0` as a *master node*.
- Join the *node* `ik8s-node-0` to the *cluster*.

Docker is already installed on both *nodes* and `apt` has been configured so that you can install the required tools.

You must use the `kubeadm` configuration file located at `/etc/kubeadm.conf` when initializing your *cluster*.

The *cluster* will be considered complete once both *nodes* are in a **Ready** state.

Question: 22/24



Question 22 ▾

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context bk8s`

Given a partially-functioning *Kubernetes cluster*, identify symptoms of failure on the *cluster*.

Determine the *node*, the failing service and take actions to bring up the failed service and restore the health of the *cluster*. Ensure that any changes are made permanently.

Hints:

- You can `ssh` to the relevant *nodes* using: `$ ssh ${NODE}` where `${NODE}` is one of **bk8s-master-0** or **bk8s-node-0**
- You can assume elevated privileges on any *node* in the *cluster* with the following command: `$ sudo -i`

Question: 23/24



Question 23▼

[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context hk8s`

Create a persistent volume with name `app-data`, of capacity `2Gi` and access mode `ReadWriteOnce`. The type of volume is `hostPath` and its location is `/srv/app-data`.

Question: 24/24



Question 24▼

[Navigate All Questions](#)

2nd Set Starts Below

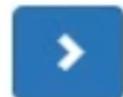
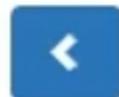
Question weight: 5%

Set configuration context: `$ kubectl config use-context k8s`

Monitor the logs of pod **foo** and:

- Extract log lines corresponding to error `file-not-found`
- Write them to `/opt/KULM00201/foo`

Question: 1/24



Question 1 ▾

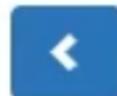
Navigate All Questions

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

List all *persistent volumes* sorted by **capacity**, saving the full `kubectl` output to `/opt/KUCC00102/volume_list`. Use `kubectl`'s own functionality for sorting the output, and do not manipulate it any further.

Question: 2/24



Question 2 ▾

[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

Ensure a single instance of pod `nginx` is running on each *node* of the *Kubernetes cluster* where `nginx` also represents the *image name* which has to be used. Do not override any *taints* currently in place.

Use **DaemonSet** to complete this task and use `ds-kusc00201` as *DaemonSet name*.

Question: 3/24



Question 3 ▾

[Navigate All Questions](#)

Question weight: 7%

Set configuration context: `$ kubectl config use-context k8s`

Perform the following tasks:

- Add an *init container* to `hungry-bear` (which has been defined in spec file `/opt/KUCC00108/pod-spec-KUCC00108.yaml`)
- The *init container* should create an empty file named `/workdir/eager.txt`
- If `/workdir/eager.txt` is not detected, the *pod* should exit
- Once the spec file has been updated with the *init container* definition, the *pod* should be created

Question: 4/24



Question 4 ▾

[Navigate All Questions](#)

English ▾

Question weight: 4%

Set configuration context: `$ kubectl config use-context k8s`

Create a pod named `kucc8` with a single app container for each of the following *images* running inside (there may be between 1 and 4 *images* specified): `nginx + redis`.

Question: 5/24



Question 5 ▾

[Navigate All Questions](#)

Question weight: 2%

Set configuration context: `$ kubectl config use-context k8s`

Schedule a pod as follows:

- Name: `nginx-kusc00101`
- Image: `nginx`
- Node selector: `disk=spinning`

Question: 6/24



Question 6 ▾

Navigate All Questions

Question weight: 4%

Set configuration context: `$ kubectl config use-context k8s`

Create a deployment as follows:

- Name: `nginx-app`
- Using container `nginx` with version `1.11.9-alpine`
- The deployment should contain `3` replicas

Next, deploy the application with new version `1.12.0-alpine`, by performing a rolling update, and record that update.

Finally, rollback that update to the previous version `1.11.9-alpine`.

Question: 7/24



Question 7 ▾

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context k8s`

Create and configure the service `front-end-service` so it's accessible through `NodePort` and routes to the existing pod named `front-end`.

Question: 8/24



Question 8 ▾

[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

Create a pod as follows:



- Name: `mongo`
- Using *image*: `mongo`
- In a new Kubernetes namespace named: `website-backend`

Question: 9/24



Question 9 ▾

[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

Create a deployment spec file that will:

- Launch 3 replicas of the `nginx` image with the label `app_runtime_stage=dev`
- `deployment` name: `kual00201`

Save a copy of this spec file to `/opt/KUAL00201/deployment_spec.yaml` (or `.json`).

When you are done, clean up (delete) any new Kubernetes API object that you produced during this task.

Question: 10/24



Question 10 ▾

[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context k8s`

Create a file: `/opt/KUCC00302/kucc00302.txt` that lists all pods that implement service **baz** in namespace **development**.

The format of the file should be one pod name per line.

Question: 11/24



Question 11 ▾

[Navigate All Questions](#)

Question weight: 9%

Set configuration context: `$ kubectl config use-context k8s`

Create a Kubernetes secret as follows:

- Name: `super-secret`
- credential: `s3kr3t`

Create a pod named `pod-secrets-via-file`, using the `redis` image, which mounts a secret named `super-secret` at `/secrets`.

Create a second pod named `pod-secrets-via-env`, using the `redis` image, which exports `credential` as `CREDENTIALS`.

Question: 12/24



Question 12 ▾

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context k8s`

Create a pod as follows:

- Name: `non-persistent-redis`
- container image: `redis`
- Volume with name: `cache-control`
- Mount path: `/data/redis`

The pod should launch in the `pre-prod` namespace and the volume must not be persistent.

Question: 13/24



Question 13 ▾

[Navigate All Questions](#)

English ▾

Question weight: 1%

Set configuration context: `$ kubectl config use-context k8s`

Scale the deployment **guestbook** to 3 pods.

Question: 14/24



Question 14 ▾

[Navigate All Questions](#)

Question weight: 2%

Set configuration context: `$ kubectl config use-context k8s`

Check to see how many *nodes* are ready (not including *nodes* tainted `NoSchedule`) and write the number to `/opt/KUCC00104/kucc00104.txt`.

Question: 15/24



Question 15 ▾

[Navigate All Questions](#)

Question weight: 2%

Set configuration context: `$ kubectl config use-context k8s`

From the *pod label name=overloaded-cpu*, find *pods* running high CPU workloads and write the name of the *pod* consuming most CPU to the file `/opt/KUTR00102/KUTR00102.txt` (which already exists).

Question: 16/24



Question 16 ▾

[Navigate All Questions](#)

Question weight: 7%

Set configuration context: `$ kubectl config use-context k8s`

Create a deployment as follows:

- Name: `nginx-dns`
- Exposed via a service `nginx-dns`
- Ensure that the service & pod are accessible via their respective DNS records
- The container(s) within any pod(s) running as a part of this deployment should use the `nginx` image

Next, use the utility `nslookup` to look up the DNS records of the service & pod and write the output to `/opt/KUNW00601/service.dns` and `/opt/KUNW00601/pod.dns` respectively. Ensure you use the `busybox:1.28` image (or earlier) for any testing, as the latest release has an upstream bug which impacts the use of `nslookup`.

Question: 17/24



Question weight: 7%

No configuration context change required for this item

Create a snapshot of the `etcd` instance running at
`https://127.0.0.1:2379`, saving the snapshot to the file path
`/data/backup/etcd-snapshot.db`.

The `etcd` instance is running `etcd` version 3.3.10.

The following *TLS* certificates/key are supplied for connecting to the server with `etcdctl`:

- CA certificate: `/opt/KUCM00302/ca.crt`
- Client certificate: `/opt/KUCM00302/etcd-client.crt`
- Client key: `/opt/KUCM00302/etcd-client.key`

Question: 18/24



Question 18 ▾

Navigate All Questions

Question weight: 4%

Set configuration context: `$ kubectl config use-context ek8s`

Set the node named `ek8s-node-0` as unavailable and reschedule all the pods running on it.

Question: 19/24



Question 19 ▾

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context wk8s`

A Kubernetes worker node, named **wk8s-node-0** is in state **NotReady**. Investigate why this is the case, and perform any appropriate steps to bring the *node* to a **Ready** state, ensuring that any changes are made permanent.

Hints:

- You can `ssh` to the failed *node* using: `$ ssh wk8s-node-0`
- You can assume elevated privileges on the *node* with the following command: `$ sudo -i`

Question: 20/24



Question 20 ▾

[Navigate All Questions](#)

Question weight: 4%

Set configuration context: `$ kubectl config use-context wk8s`

Configure the `kubelet` `systemd`-managed service, on the *node* labelled with `name=wk8s-node-1`, to launch a *pod* containing a single *container* of *image jenkins* named `webtool` automatically. Any spec files required should be placed in the `/etc/kubernetes/manifests` directory on the *node*.

Hints:

- You can `ssh` to the appropriate *node* using: `$ ssh wk8s-node-1`
- You can assume elevated privileges on the *node* with the following command: `$ sudo -i`

Question: 21/24



Question 21 ▾

[Navigate All Questions](#)

Important: For this item, you will have to `ssh` to the *nodes* `ik8s-master-0` and `ik8s-node-0` and complete all tasks on these *nodes*. Ensure that you return to the *base node* (hostname: `node-1`) when you have completed this item.

As an administrator of a small development team, you have been asked to set up a *Kubernetes cluster* to test the viability of a new *application*.

You must use `kubeadm` to perform this task. Any `kubeadm` invocations will require the use of the `--ignore-preflight-errors=all` option.

- Configure the *node* `ik8s-master-0` as a *master node*.
- Join the *node* `ik8s-node-0` to the *cluster*.

Docker is already installed on both *nodes* and `apt` has been configured so that you can install the required tools.

You must use the `kubeadm` configuration file located at `/etc/kubeadm.conf` when initializing your *cluster*.

The *cluster* will be considered complete once both *nodes* are in a **Ready** state.

Question: 22/24

Question weight: 4%

Set configuration context: `$ kubectl config use-context bk8s`

Given a partially-functioning *Kubernetes cluster*, identify symptoms of failure on the *cluster*.

Determine the *node*, the failing service, and take actions to bring up the failed service and restore the health of the *cluster*. Ensure that any changes are made permanently.

Hints:

- You can `ssh` to the relevant *nodes* using: `$ ssh ${NODE}` where `${NODE}` is one of **bk8s-master-0** or **bk8s-node-0**
- You can assume elevated privileges on any *node* in the *cluster* with the following command: `$ sudo -i`

Question: 23/24



Question 23 ▾

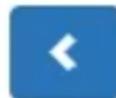
[Navigate All Questions](#)

Question weight: 3%

Set configuration context: `$ kubectl config use-context hk8s`

Create a *persistent volume* with name `app-data`, of capacity `1Gi` and access mode `ReadWriteMany`. The type of volume is `hostPath` and its location is `/srv/app-data`.

Question: 24/24



Question 24 ▾

[Navigate All Questions](#)

3rd Set Starts Below

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Monitor the logs of pod bar and:

- Extract log lines corresponding to error file-not-found
- Write them to /opt/KULM00201/bar

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Monitor the logs of pod bar and:

- Extract log lines corresponding to error file-not-found
- Write them to /opt/KULM00201/bar



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

List all persistent volumes sorted by capacity , saving the full kubectl output to /opt/KUCC00102/pv_list . Use kubectl 's own functionality for sorting the output, and do not manipulate it any further.



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Perform the following tasks:

- Add an init container to `lumpy-koala` (which has been defined in spec file `/opt/KUCC00108/pod-spec-KUCC00108.yaml`)
 - The init container should create an empty file named `/workdir/faithful.txt`
 - If `/workdir/faithful.txt` is not

Perform the following tasks:

- Add an init container to `lumpy-koala` (which has been defined in spec file
`/opt/KUCC00108/pod-spec-KUCC00108.yaml`
)
- The init container should create an empty file named `/workdir/faithful.txt`
- If `/workdir/faithful.txt` is not detected, the pod should exit
- Once the spec file has been updated with the init container definition, the pod should be created



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Ensure a single instance of pod `nginx` is running on each node of the Kubernetes cluster where `nginx` also represents the image name which has to be used. Do not override any taints currently in place.

Use **DaemonSet** to complete this task and use `ds-kusc00201` as DaemonSet name.

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create a pod named `kucc1` with a single app container for each of the following images running inside (there may be between 1 and 4 images specified):
`nginx + redis`.



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Schedule a pod as follows:

- Name: nginx-kusc00101
- Image: nginx
- Node selector: disk=spinning



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create a deployment as follows:

- Name: nginx-app
- Using container nginx with version 1.11.9-alpine
- The deployment should contain 3 replicas

Next, deploy the application with new version 1.13.0-alpine , by performing a rolling update.

Finally, rollback that update to the

Task

Create a deployment as follows:

- Name: `nginx-app`
- Using container `nginx` with version `1.11.9-alpine`
- The deployment should contain 3 replicas

Next, deploy the application with new version `1.13.0-alpine`, by performing a rolling update.

Finally, rollback that update to the previous version `1.11.9-alpine`.



Flag this to ...

I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create and configure the service `front-end-service` so it's accessible through `NodePort` and routes to the existing pod named `front-end`.



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create a pod as follows:

- Name: nginx
- Using image: nginx
- In a new Kubernetes namespace named: website-backend

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create a deployment spec file that will:

- Launch 4 replicas of the redis image with the label `app_runtime_mode=prod`
- deployment name: `kual00201`

Save a copy of this spec file to

`/opt/KUAL00201/deploy_spec.yaml` (or
`/opt/KUAL00201/deploy_spec.json`).

When you are done, clean up (delete)

```
[student@node-1] $ | kubectl config use-context k8s
```

Task

Create a deployment spec file that will:

- Launch 4 replicas of the redis image with the label app_runtime_mode=prod
- deployment name: kual00201

Save a copy of this spec file to /opt/KUAL00201/deploy_spec.yaml (or /opt/KUAL00201/deploy_spec.json).

When you are done, clean up (delete) any new Kubernetes API object that you produced during this task.

Flag this to ...

I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create a file:

/opt/KUCC00302/kucc00302.txt that lists all pods that implement service bar in namespace development .

The format of the file should be one pod name per line.



Flag this to ...

I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create a Kubernetes secret as follows:

- Name: super-secret
- password : s3kr3t

Create a pod named pod-secrets-via-file , using the redis image, which mounts a secret named super-secret at /secrets .

Create a second pod named pod-secrets-via-env , using the redis image, which exports password as TOPSECRET .

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create a pod as follows:

- Name: non-persistent-redis
- container image: redis
- Volume with name: cache-control
- Mount path: /data/redis

The pod should launch in the pre-prod namespace and the volume **must not** be persistent.

Task weight: 1%

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Scale the deployment `loadbalancer` to 3 pods.



Flag this to ...

I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Check to see how many worker nodes are ready (not including nodes tainted NoSchedule) and write the number to /opt/KUCC00104/kucc00104.txt .

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

From the pod label

name=overloaded-cpu , find pods running high CPU workloads and write the name of the pod consuming most CPU to the file

/opt/KUTR00102/KUTR00102.txt (which already exists).

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context k8s
```

Task

Create a deployment as follows:

- Name: nginx-random
- Exposed via a service
`nginx-random`
- Ensure that the service & pod are accessible via their respective DNS records
- The container(s) within any pod(s) running as a part of this deployment should use the `nginx` image

Create a deployment as follows:

- Name: `nginx-random`
- Exposed via a service
`nginx-random`
- Ensure that the service & pod are accessible via their respective DNS records
- The container(s) within any pod(s) running as a part of this deployment should use the `nginx` image

Next, use the utility `nslookup` to look up the DNS records of the service & pod and write the output to
`/opt/KUNW00601/service.dns` and
`/opt/KUNW00601/pod.dns` respectively.

No configuration context
change required for this item.



Task

Create a snapshot of the `etcd` instance running at <https://127.0.0.1:2379>, saving the snapshot to the file path `/data/backup/etcd-snapshot.db`.

The following TLS certificates/key are supplied for connecting to the server with `etcdctl`:



- CA certificate:
`/opt/KUCM00302/ca.crt`
- Client certificate:
`/opt/KUCM00302/etcd-client.`

saving the snapshot to the file path
`/data/backup/etcd-snapshot.db` .

The following TLS certificates/key are supplied for connecting to the server with `etcdctl` :

- CA certificate:
`/opt/KUCM00302/ca.crt`
- Client certificate:
`/opt/KUCM00302/etcd-client.crt`
- Client key:
`/opt/KUCM00302/etcd-client.key`

 Flag this to ...

I am satisfied, ...

Task weight: 4%

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context ek8s
```

Task

Set the node named `ek8s-node-1` as unavailable and reschedule all the pods running on it.



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context wk8s
```

Task

A Kubernetes worker node, named `wk8s-node-0` is in state `NotReady`. Investigate why this is the case, and perform any appropriate steps to bring the node to a `Ready` state, ensuring that any changes are made permanent.

You can `ssh` to the failed node using:



```
[student@node-1] $ | ssh wk8
```

perform any appropriate steps to bring the node to a **Ready** state, ensuring that any changes are made permanent.

You can **ssh** to the failed node using:

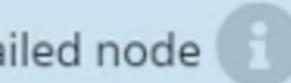
```
[student@node-1] $ | ssh wk8  
s-node-0
```

You can assume elevated privileges on the node with the following command:

```
[student@wk8s-node-0] $ | su  
do -i
```



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context wk8s
```

Task

Configure the kubelet systemd-managed service, on the node labelled with name=wk8s-node-1 , to launch a pod containing a single container of image jenkins named app automatically. Any spec files required should be placed in the /etc/kubernetes/manifests directory on the node.

You can ssh to the appropriate node using:



the `/etc/kubernetes/manifests` directory on the node.

You can ssh to the appropriate node using:

```
[student@node-1] $ | ssh wk8  
s-node-1
```

You can assume elevated privileges on the node with the following command:

```
[student@wk8s-node-1] $ | su  
do -i
```

Flag this to ...

I am satisfied, ...

No configuration context
change required for this item.



For this item, you will have to **ssh** to the nodes **ik8s-master-0** and **ik8s-node-0** and complete all tasks on these nodes. Ensure that you return to the base node (hostname: **node-1**) when you have completed this item.



Context

As an administrator of a small development team, you have been asked to set up a Kubernetes cluster to test the viability of a new application.

Context

As an administrator of a small development team, you have been asked to set up a Kubernetes cluster to test the viability of a new application.

Task

You must use `kubeadm` to perform this task. Any `kubeadm` invocations will require the use of the `--ignore-preflight-errors=all` option.

- Configure the node `ik8s-master-0` as a master node.
- Join the node `ik8s-node-0` to the cluster.

You must use the `kubeadm` configuration file located at `/etc/kubeadm.conf` when initializing

- Configure the node `ik8s-master-0` as a master node.
- Join the node `ik8s-node-0` to the cluster.

You must use the `kubeadm` configuration file located at `/etc/kubeadm.conf` when initializing your cluster.

The cluster will be considered complete once both nodes are in a **Ready** state.

Docker is already installed on both nodes and `apt` has been configured so that you can install the required tools.



Flag this to ...

I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context bk8s
```

Task

Given a partially-functioning Kubernetes cluster, identify symptoms of failure on the cluster.

Determine the node, the failing service, and take actions to bring up the failed service and restore the health of the cluster. Ensure that any changes are made permanently.

You can ssh to the relevant



cluster, identify symptoms or failure on the cluster.

Determine the node, the failing service, and take actions to bring up the failed service and restore the health of the cluster. Ensure that any changes are made permanently.

You can ssh to the relevant nodes (bk8s-master-0 or bk8s-node-0) using:

```
[student@node-1] $ | ssh <nodename>
```

You can assume elevated privileges on any node in the cluster with the following command:

```
[student@nodename] $ | sudo -i
```

made permanently.



You can ssh to the relevant nodes (bk8s-master-0 or bk8s-node-0) using:

```
[student@node-1] $ | ssh <no  
dename>
```

You can assume elevated privileges on any node in the cluster with the following command:

```
[student@node-1] $ | sudo  
-i
```



Flag this to ...



I am satisfied, ...

Set configuration context:



```
[student@node-1] $ | kubectl  
config use-context hk8s
```

Task

Create a persistent volume with name `app-data`, of capacity `1Gi` and access mode `ReadWriteOnce`. The type of volume is `hostPath` and its location is `/srv/app-data`.



Flag this to ...